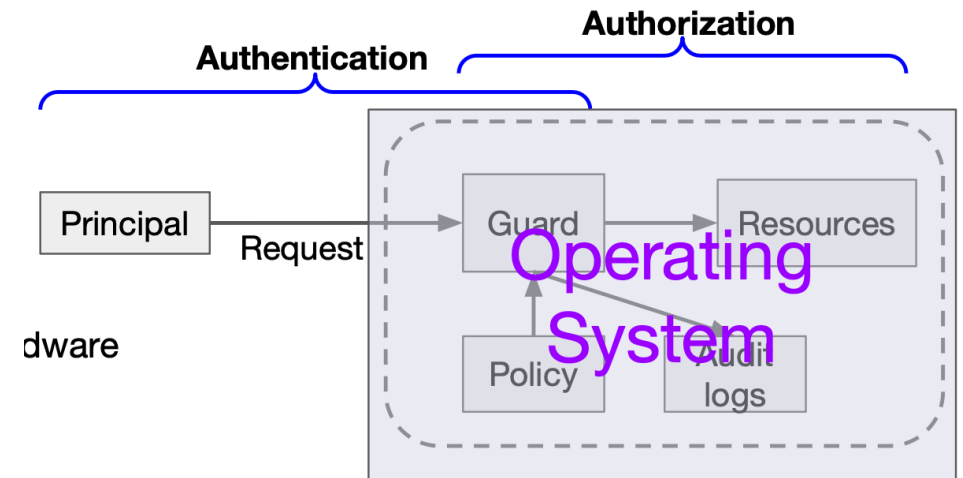
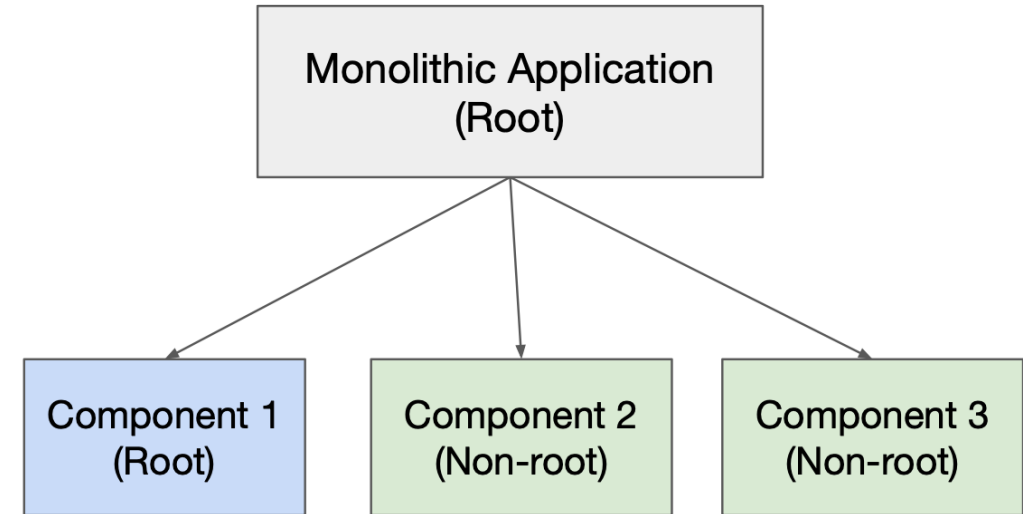


Isolation

Dileepa Fernando

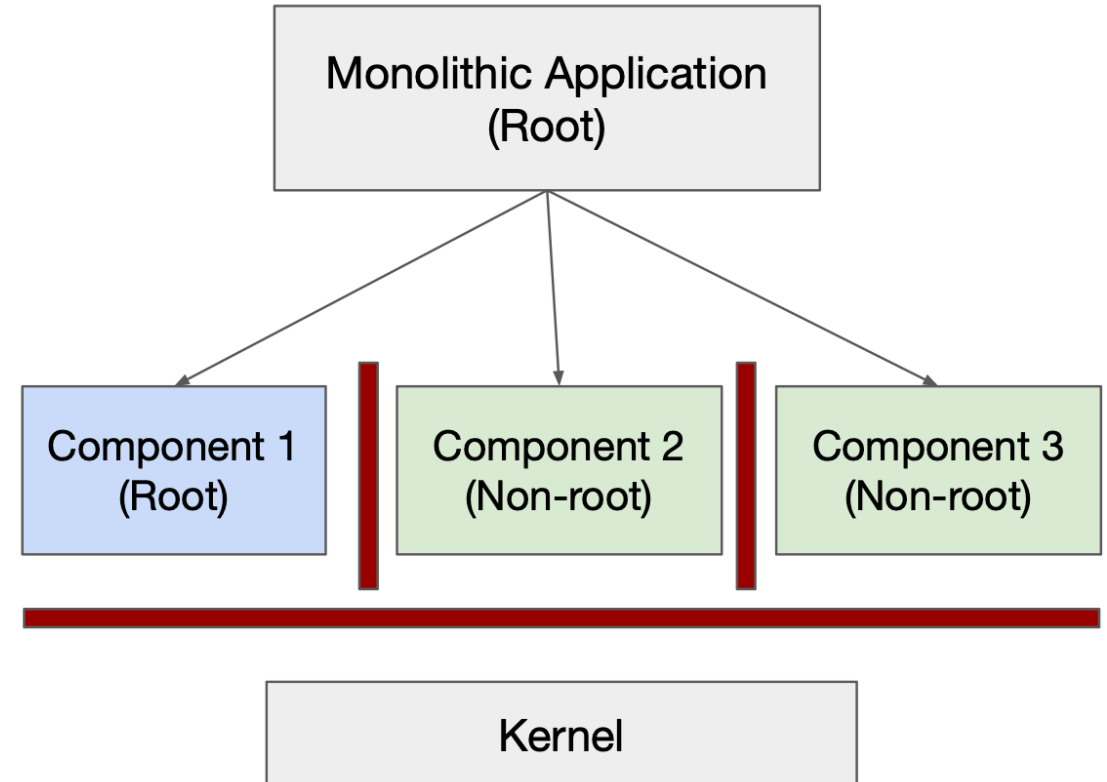
Story (Privilege Separation)

- Modularize Monolithic Applications
 - Different folders?
 - *chroot*?
- For each module
 - Escalate privilege (setuid-binary, sudo, ..)
 - Sensitive operations
 - Drop privilege (setuid(u))
 - After sensitive operations
 - Least privilege principle
- Assumptions
 - We know what privilege to assign
 - OS enforces




Isolation Needed for Priv. Separation

- The OS isolated
 - The processes
- A cannot affect B directly
- Errors are contained
- 100% Isolation is not desirable
- Controlled communication needed




Isolation (General)

- Lower Level Enforces
- TCB is the lower level
- Example:
 - App separation?
 - OS separation?
 - Hypervisor separation?



Isolation (General)

- Different Isolation levels
- Different security requirements
 - Keep Attackers out
 - Keep attackers in
- Different performance requirements
 - Container vs VM vs in app




Overview

- Introduction
- Kernel Isolation
- [Application Isolation – Virtual Machine](#)
- [Application Isolation - Container](#)

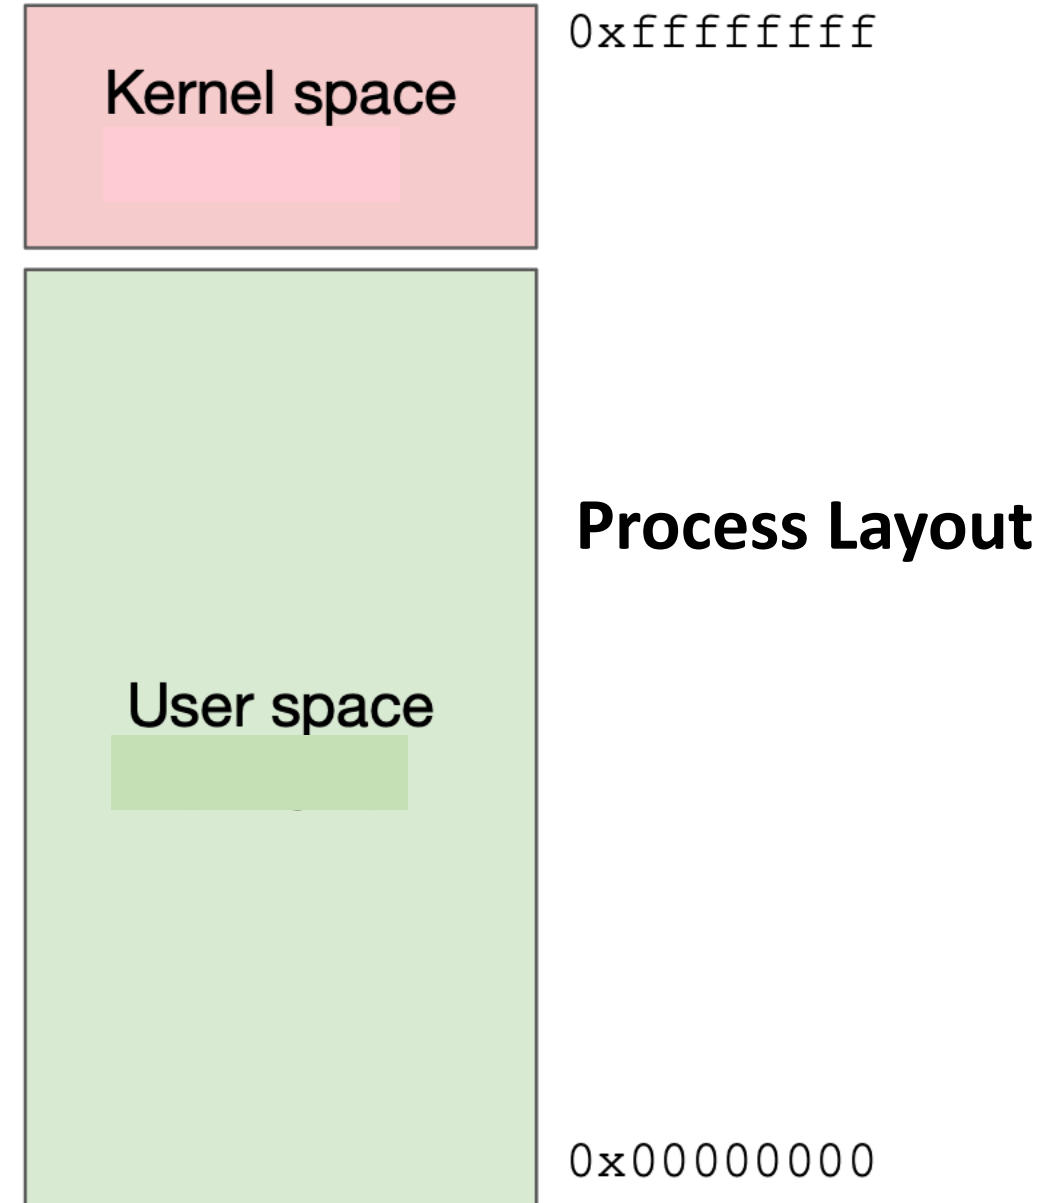
Kernel Isolation

- Role
 - Manage resources
 - Abstract resources
- Security Requirements
 - [Isolate itself](#)
 - [Isolate processes among each other](#)
- TCB
 - CPU + Kernel Code
- Threat Model
 - Applications
 - Administrator
 - Hardware vendor?
 - Physical Tampering?



Kernel Isolating Itself

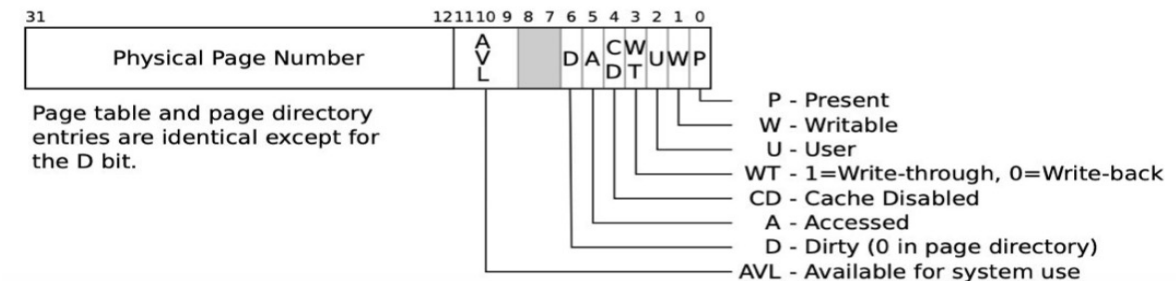
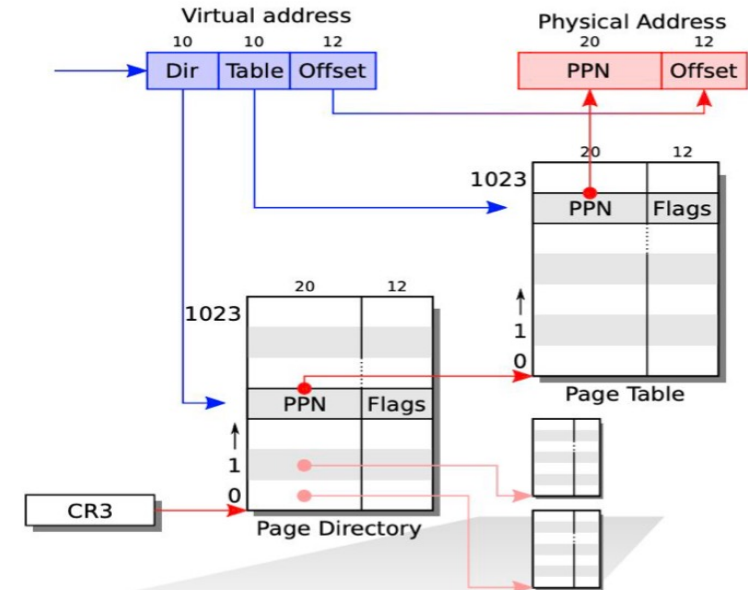

- [Memory Isolation \(Virtual\)](#)
 - In 32 bit, 1GB dedicated for kernel
 - User process cannot access kernel space Directly
- [Instruction Isolation](#)
 - CPU runs in different privilege levels
 - Instructions are restricted at some levels
 - Not to be confused with **process privilege**
 - But idea is similar



Kernel Isolating Itself

- **Memory Isolation**

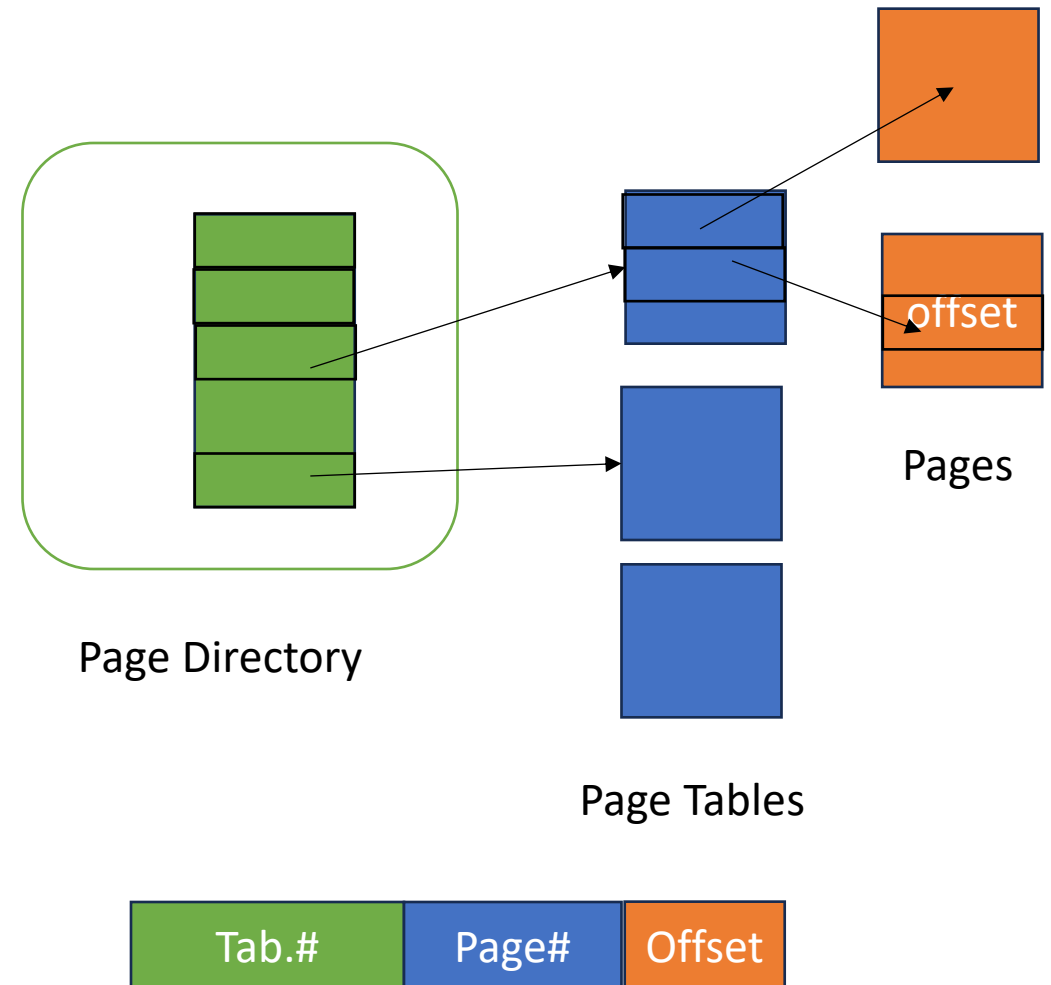

- Virtual -> Physical Translation
 - Page tables reside in kernel
 - Referred from CPU (CR3 register)
- Cannot write directly to physical mem.



Kernel Isolating Itself

- **Memory Isolation**


- Virtual -> Physical Translation
- Happens inside CPU
- User never knows



Kernel Isolating Itself

- **Instruction Isolation**

- CPU executes instructions according to privilege levels (Ring)
- Ring 0 (Highest Privilege)
 - Any instruction executed
- Ring 3 (Lowest Privilege)
 - Any user process (including uid 0)
 - **Root user can attack user resources**
 - But not the OS



Kernel Isolating Itself


- **Instruction Isolation**

- Saved in Code Segment Register
- *info reg*
- CPL checked for each instruction
 - What is CPL?



CPL = 0: ring 0


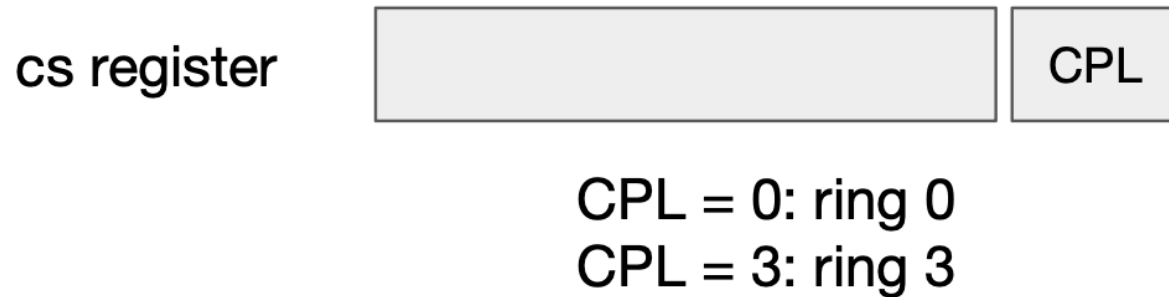
CPL = 3: ring 3



Kernel Isolating Itself


- **Instruction Isolation**

- Ring 0 at boot
- Only changes with interrupts
 - System calls
 - Hardware interrupts



Kernel Isolating Itself



- General Isolation Idea
 - Kernel/ User Space separation
 - User space process -> Ring 0
 - CPL=0
 - Kernel space process -> Ring 3
 - CPL=3
- Data Isolation
 - Kernel Space data access prohibited for Ring 3
 - MOV EAX [0xffff~~ff~~]
- Code Isolation
 - Kernel instructions are prohibited for Ring 3
 - C~~xx~~



Kernel Isolating Itself


- Data Isolation
 - Read/Write kernel space not allowed for Ring 3
 - Can we jump to a kernel address?
 - JMP ADDR
 - Can we overwrite page table directly?
 - Can we point CR3 to somewhere else?

Page Table Entry (U Flag)




Kernel Isolating Itself

- We cannot jump to a kernel address
- But a user need to access kernel space
- Need safe transition
 - Well defined exit/entry/behavior
 - Safe change of privilege level
 - **Interrupt**
 - Stop normal execution
 - Transition to kernel
 - Software – System Calls
 - Hardware – Timer, Faults



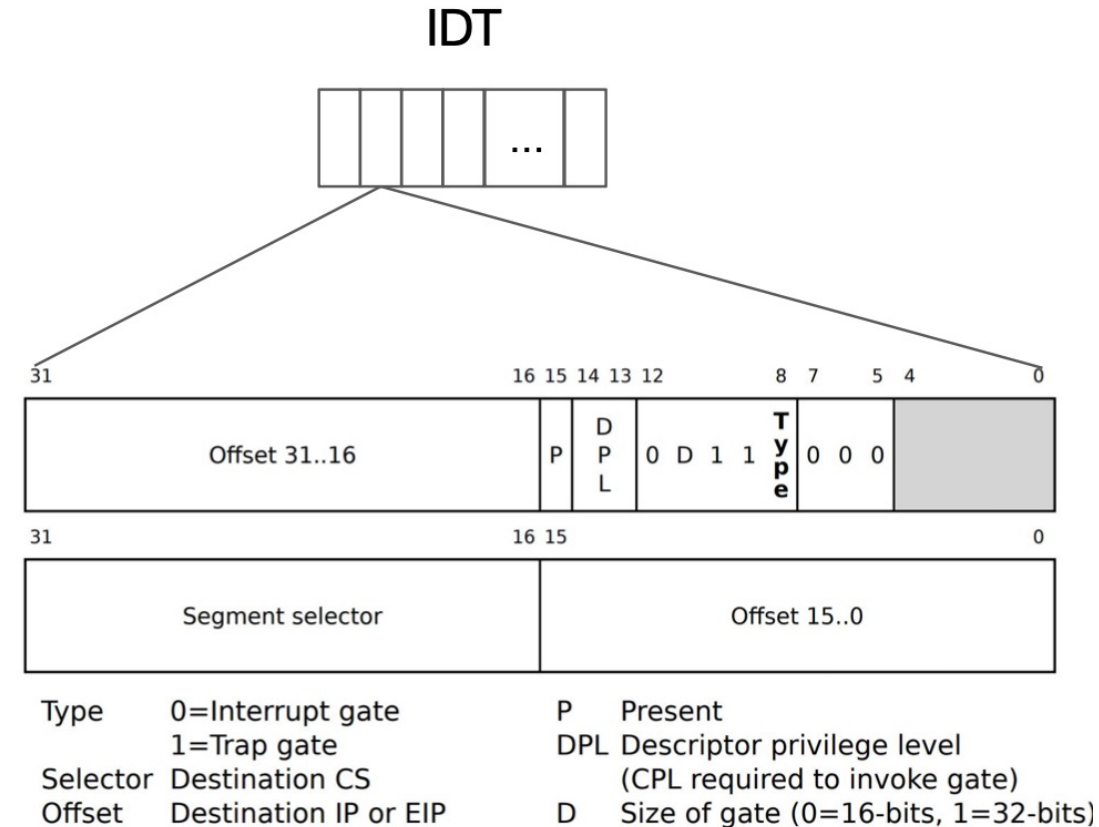
Kernel Isolating Itself

- We cannot jump to a kernel address
- But a user need to access kernel space
- Need safe transition
 - INT N command
 - Control to IDT[N] (kernel space)
 - SET CPL
 - IRET
 - SET CPL
 - Return to user space
 - Security Principle?



Kernel Isolating Itself

- Interrupt Descriptor Table
 - 256 Entries
 - Each Entry 64 bytes
 - Address located in %idtr reg
- When INT N received (Detail)
 - Control to IDT[N]
 - Check $CPL \leq DPL$
 - SET CS, EIP
 - Interrupt Service Routine
 - IRET
 - SET CS, EIP
 - Return



Kernel Isolating Itself




- Question (Comprehension)
 - If SET CPL=0, JMP <addr> is allowed
 - What can happen?
 - Setting CPL in kernel space
- If JMP <a kernel addr> SET CPL=0 is allowed
- What can happen?
- Not all kernel addresses should be jumped to
- Changing Privilege level only in Kernel space
- Only jumping to well defined Kernel functions

Kernel Isolating Itself

- Example (System Call)

- INT 0x80

- IDT entry
- DPL 3
- Address of ISR code
 - 0xffff000a
- Final value of CPL
 - 0x00000000





Kernel Isolating Itself

- Example (Page Fault)

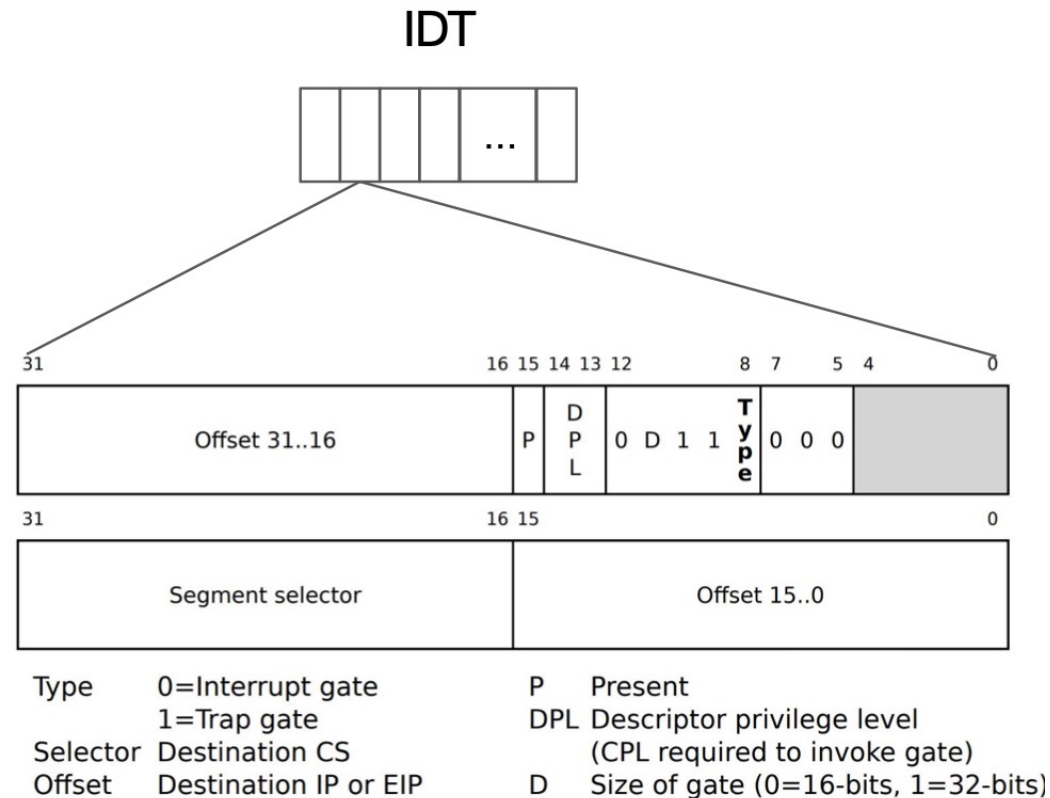
- INT 0xe

- IDT entry
- DPL 0
- Address of ISR code
 - 0xffff001a
- Final value of CPL
 - 0x00000000



Kernel Isolating Itself

- Question
 - How does hardware stop you if you called INT 0xe?

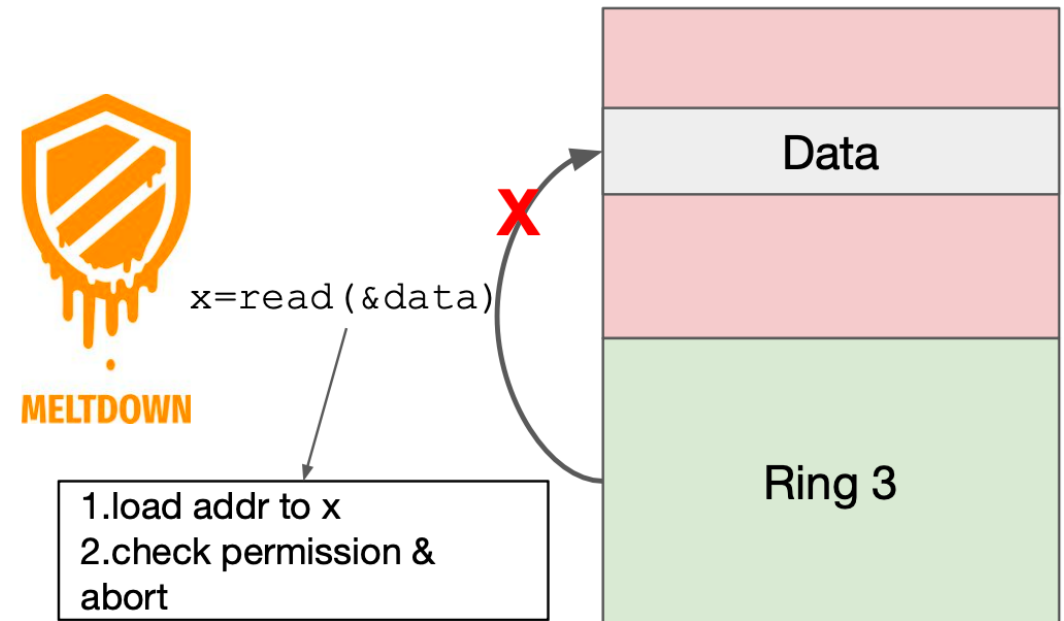


Kernel Isolating Itself

- Summary
 - Data Isolation
 - Virtually Separate kernel and user space
 - Physical separation with U flag
 - User cannot directly access kernel space (MOV, SET, JMP)
 - Instruction Isolation
 - Any instruction in kernel space
 - Restricted instructions in user space
 - SET CR3, SET CS etc. restricted
 - Safe transition (Controlled Interaction)
 - INT 0x80 from user space
 - INT 0xe from hardware
 - SET CPL only in kernel space

Kernel Isolating Itself

- Can we afford to write sloppy kernel code?
- Meltdown attack
 - Does this work?
 - Works when CPU speculates
 - Cache before check CPL/ U flag




Isolation Among Processes (Separation)

- Each Process has their own Virtual Address Space

- May overlap some times

- 0xdeadbeef in Proc. A
 - 0xdeadbeef in Proc. B




- Their Physical Address spaces do not overlap

- 0xdeadbeef in Proc. A and Proc. B map to different physical addresses
 - Except for explicit shared memory
 - **Enforced by Kernel**

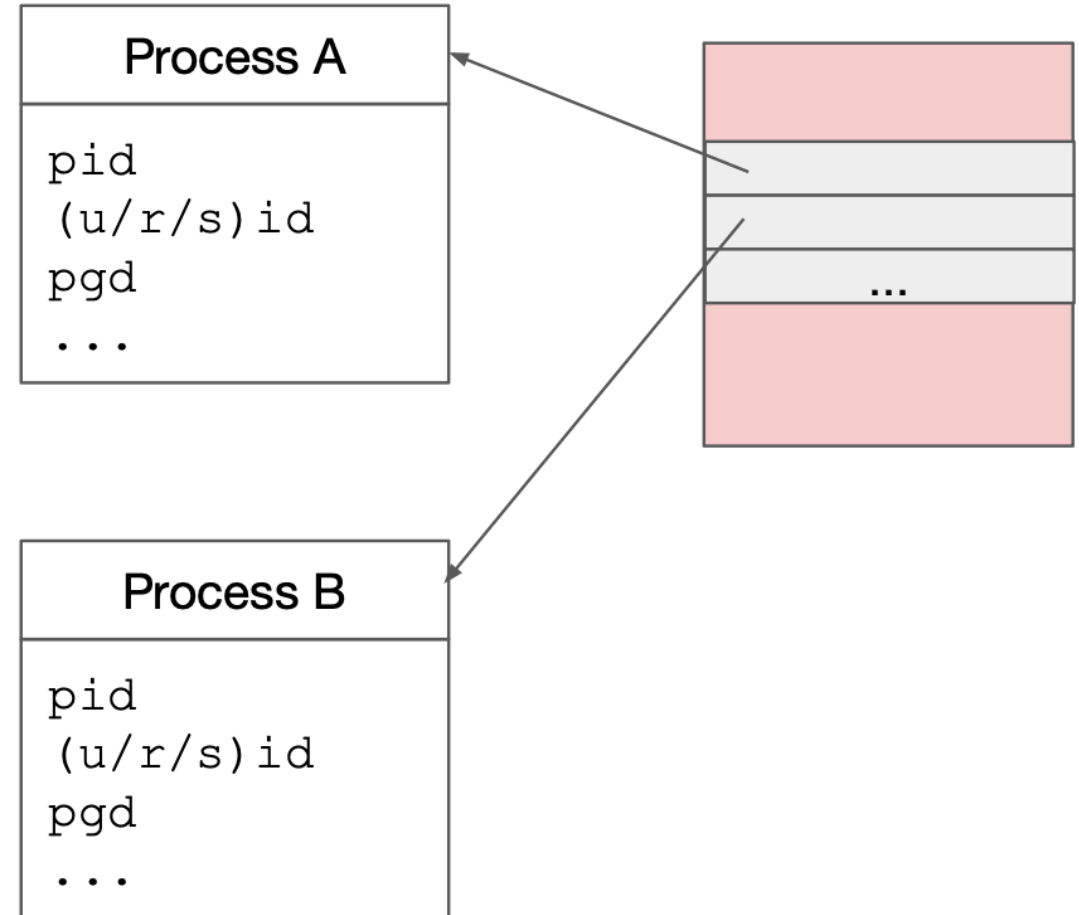
Isolation Among Processes (Enforcement)

- Can Proc. A access page table of Proc. B?
 - When Proc. A running
 - CR3 is page directory of A
 - It has to be changed
 - To access a different page table
 - CR3 cannot be changed from user space
- Can we have our own page table?



Isolation Among Processes

- Kernel Ensures no overlap
 - Ex: For Malloc
- During Context Switch
 - Kernel takes control
 - SET CR3 to proc->pgd
 - Flush TLB



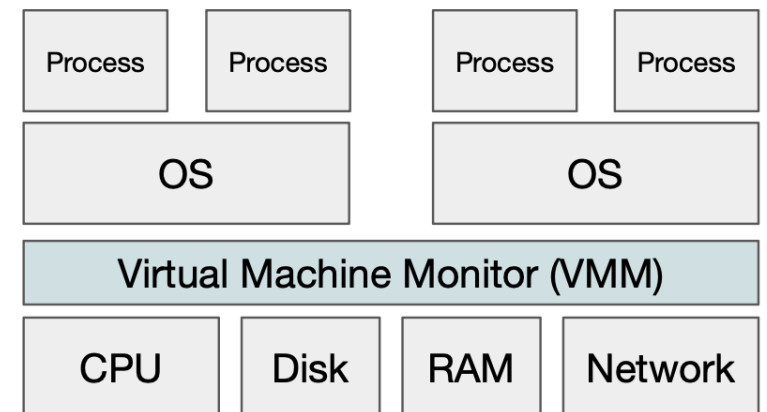


Kernel Isolation

- Summary
 - CPU privilege levels and safe transitions
 - User/ Kernel space separation (Virtual Memory)
 - Separation between Processes (Virtual Memory)
 - Enforcement by Kernel and CPU together
 - Kernel, CPU vulnerabilities lead to attacks



Application Isolation – Virtual Machine

- Old Idea – Popek & Goldberg 1974
 - Host and Guest concept
 - Real Hw (Host Hw)
 - **Simulated** Hw (VMM, hypervisor)
 - First Sw contact with **simulated** Hw (Guest OS)
 - We focus on VMM directly running on Hw
 - Why do we call VMM the hypervisor?



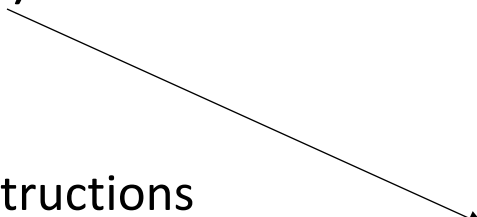
Application Isolation – Virtual Machine

- Simulation

- Assume you write a program (Virtual Machine)
 - Accepts a binary as input (Ex: ./demo)
 - Keep data structures for CPU, Memory, etc.
 - Update the data structures according to binary instructions

- Virtual Machine Monitor

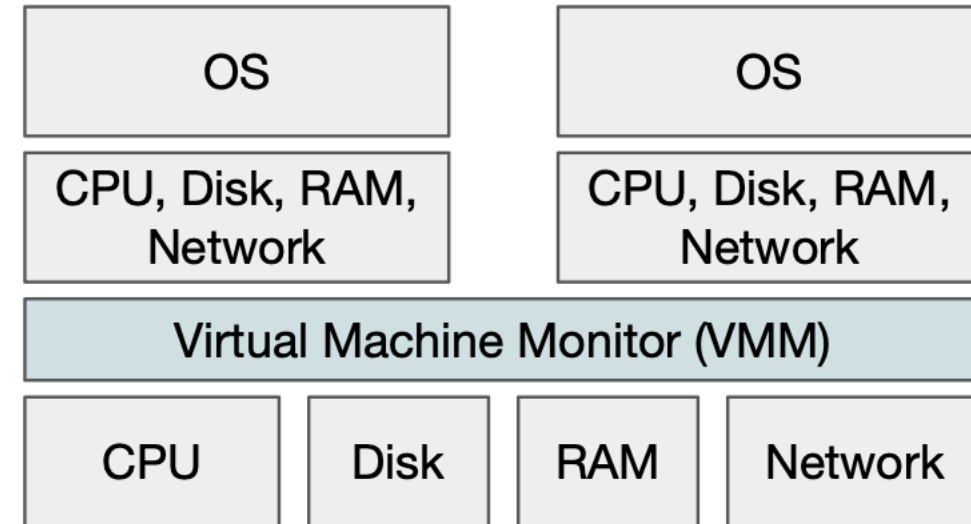
- Manages/ Runs virtual machines
- Efficiently distribute Hw resources



```
while ( ip < code.length ) {  
    int opcode = code[ip]; // fetch  
    if ( trace ) System.err.printf("%-35s", disInstr());  
    ip++;  
    switch (opcode) {  
        case ICONST:  
            int v = code[ip];  
            ip++;  
            sp++;  
            stack[sp] = v;  
            break;  
        case PRINT:  
            v = stack[sp];  
            sp--;  
            System.out.println(v);  
            break;  
        case GLOAD :  
            int addr = code[ip];  
            ip++;  
    }  
}
```

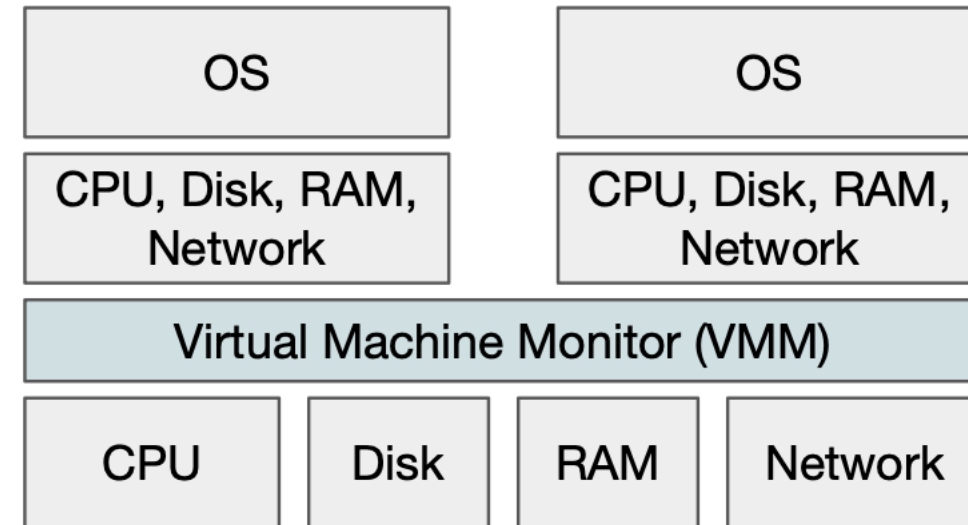
Application Isolation – Virtual Machine

- To VMM
 - Guest OS is a user program
 - Guest OS has its own virtual memory space
- To guest OS
 - It sees the Hw simulation as real
 - Virtual Hw executes guest OS



Application Isolation – Virtual Machine


- The Virtual Machine Monitor
- Requirement 1: Protection
 - VMM protects itself
 - Guest OS kept inside simulation
- Requirement 2: Illusion
 - Guest OS must not realize the simulation
 - No check inside guest OS should reveal this
- Requirement 3: Performance
- Two more security principles
 - Defense in-depth
 - Small TCB




Application Isolation – Virtual Machine

- Approach (Trap and Emulate)
 - VMM runs in Ring 0
 - Guest OS/ processes run in Ring 3
 - CPU executes Guest Instruction
 - ADD, XOR, PUSH etc.
- Privileged Instructions
 - CLI, LCR3
 - CPU raises exceptions
 - Handled by VMM

add %eax, %ebx



cli



Application Isolation – Virtual Machine

- Approach (Trap and Emulate) – Virtualize Memory

- Guest

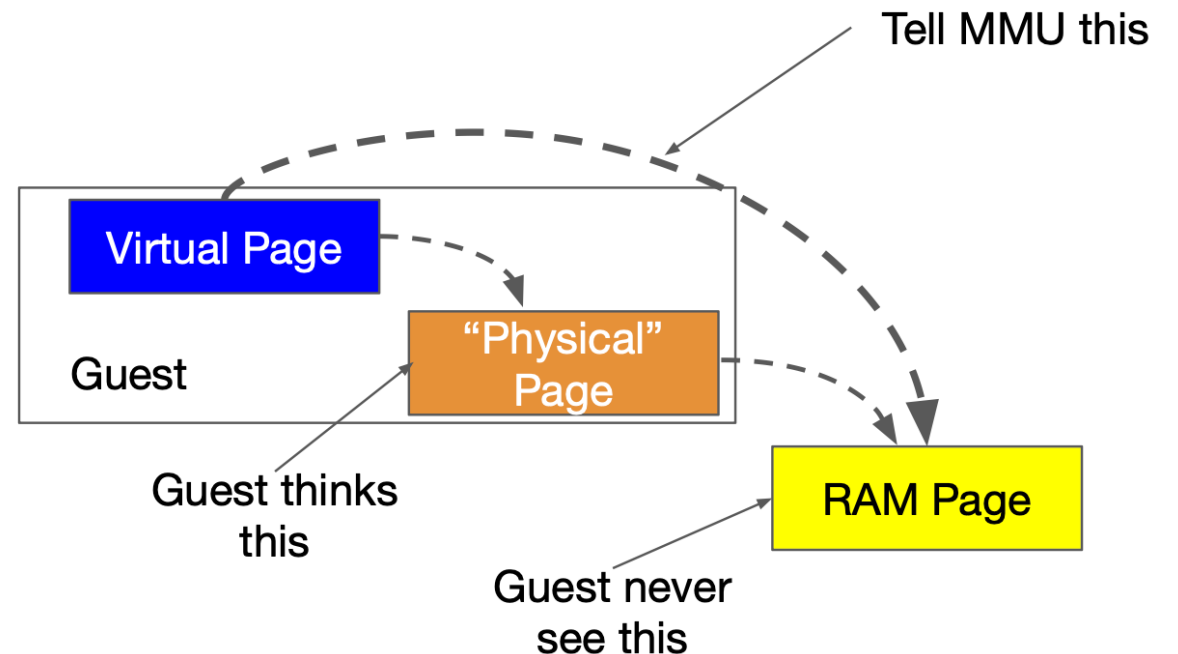
- Guest Virtual
 - Guest Physical

- Host

- Guest Physical
 - Host Physical


- Modern Implementations

- Guest Virtual
 - Host Physical
 - Shadow page table



Application Isolation – Virtual Machine

- Approach (Trap and Emulate)
 - Shadow Page Table
 - One per Guest Process
 - Only VMM sees



Application Isolation – Virtual Machine

- Does Trap and Emulate Meet Requirements?
 - Protection, Illusion, Performance
- Alternative 1: Binary Translation (VM Ware)
 - Translate Troublesome instructions to safer ones
 - `MOV CS EAX -> INT XX`
- Alternative 2: Paravirtualization (Xen)
 - VM knows it is in simulation
 - Guest OS is modified to communicate with VMM
- Alternative 3: Hardware Support (Intel-VTX, AMD-SVM)
 - Privilege handling
 - Memory/IO management

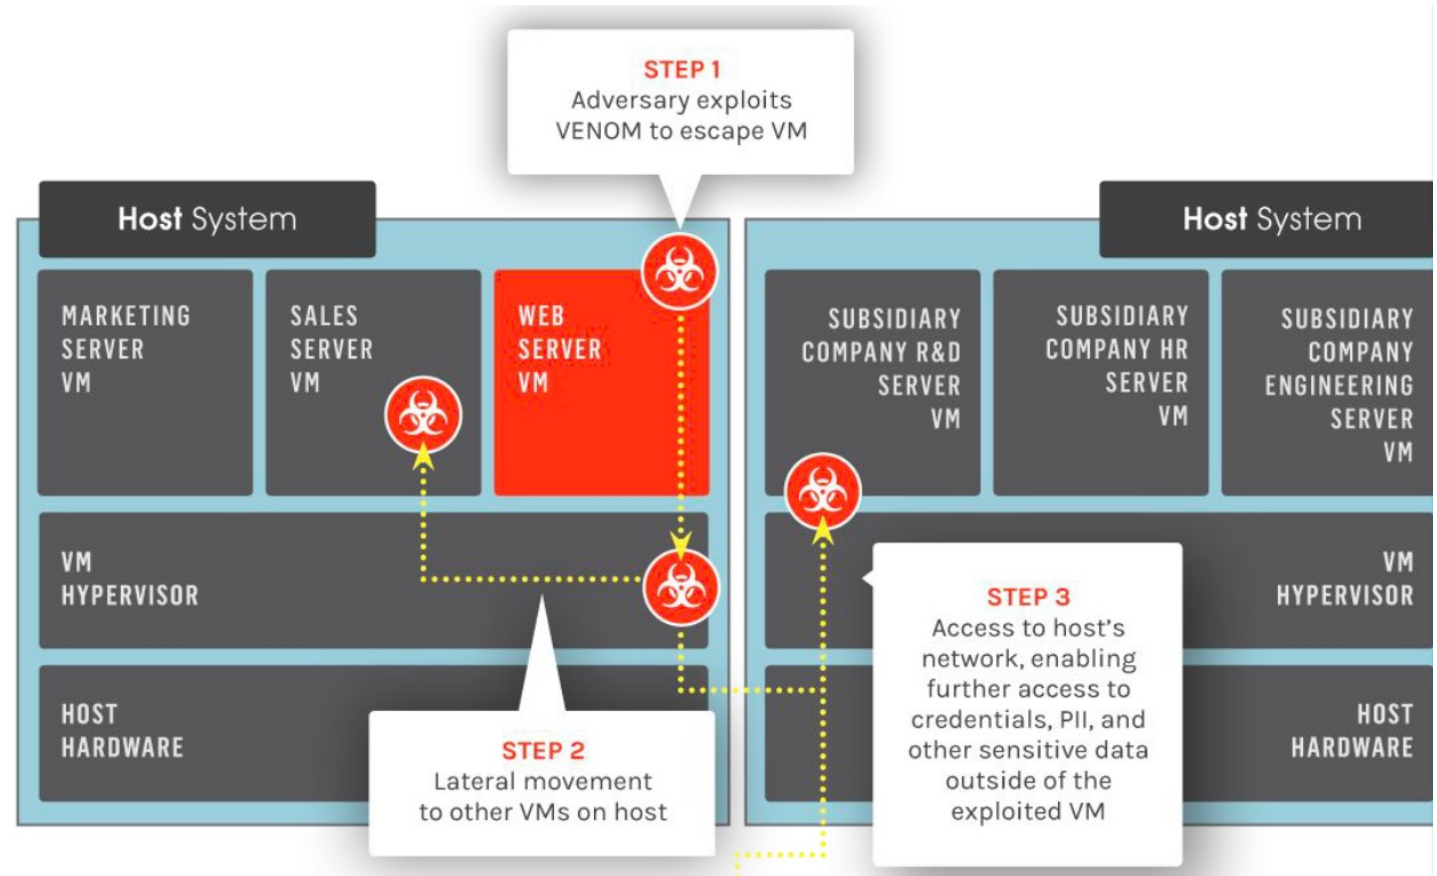
Application Isolation – Virtual Machine

- Summary

- Simulation good for Isolation
- Another Abstraction Layer
- Fast Enough?




- Attack?

- Buffer Overflow in VMM
 - VENOM




Application Isolation – Containers

- Isolation So Far
 - TCB – Hw, Kernel
 - Virtualized Hardware
 - Performance, Isolation Trade-off
- Container – Somewhere in the middle
- Virtualize OS




Application Isolation - Containers

- Idea1: Virtualize and Isolate Kernel Resources
- Idea2: Limit Resources for each container
- Idea3: Limit Kernel Access for each container (System Call Filter)
- All available in linux distribution




Application Isolation - Containers

- **Idea1: Virtualize and Isolate Kernel Resources**
- Resources: Name Spaces
 - Process A in name space X != Process A in name space Y
 - *man namespace*
- Analogies
 - Path in Chroot
 - Address A in two virtual mem. spaces



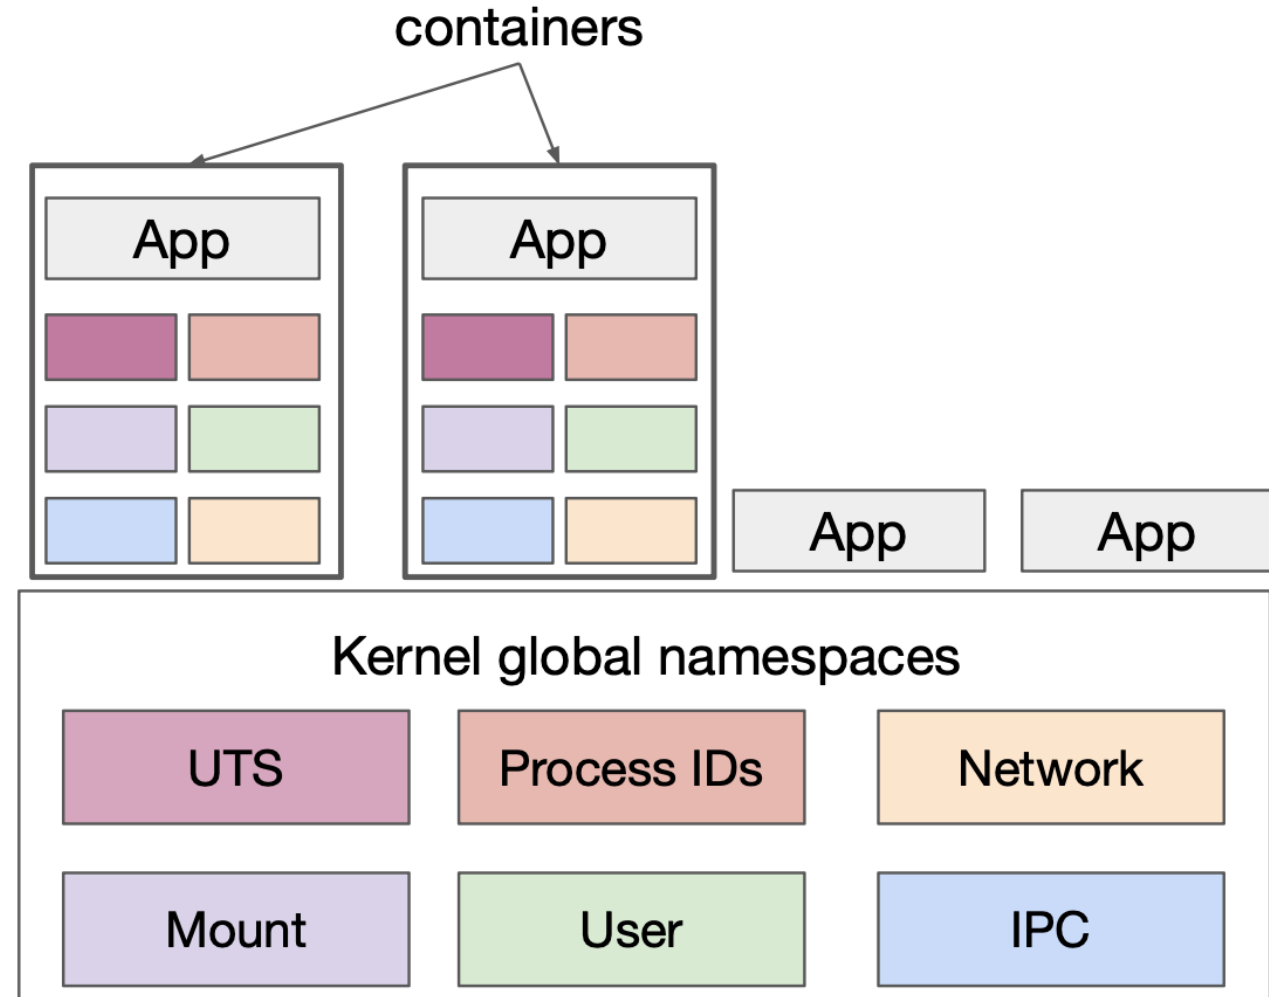
Application Isolation - Containers

- **Idea2: Limit Resources for each container (CPU, RAM, IO)**
- cgroups
 - Set Resource limit to a group of processes
 - *man cgroups*
- namespaces + cgroups
 - Start a new process in namespace
 - Add it and children to a new cgroup



Application Isolation - Containers

- **Idea3: Limit Kernel Access for each container (System Call Filter)**
- Linux has >300 system calls
- Restrict accessible system calls to container
- Goal: Reduce TCB
- Seccom:
 - A file of black/white list
 - Include the file as running option



Application Isolation - Containers

- **Idea3: Limit Kernel Access for each container (System Call Filter)**
- Linux has >300 system calls
- Restrict accessible system calls to container
- Goal: Reduce TCB
- Seccom:
 - Pass a filter to the kernel
 - Process can add filters

```
"defaultAction": "SCMP_ACT_ALLOW",
"architectures": [
  "SCMP_ARCH_X86_64",
  "SCMP_ARCH_X86",
  "SCMP_ARCH_X32"
],
"syscalls": [
  {
    "name": "chmod",
    "action": "SCMP_ACT_ERRNO",
    "args": []
  },
  {
    "name": "fchmod",
    "action": "SCMP_ACT_ERRNO",
    "args": []
  },
  {
    "name": "fchmodat",
    "action": "SCMP_ACT_ERRNO",
    "args": []
  },
  {
    "name": "chown",
    "action": "SCMP_ACT_ERRNO",
    "args": []
  }
]
```

Application Isolation - Containers

- Namespace + cgroups + seccomp = LXC
- LXC + Management Tools = Docker

- Docker vs chroot?
- Docker vs VMM?

Summary

- Isolation necessary for security
- Isolation at different levels
 - Kernel/ Hypervisor
 - Process
 - VM
 - Container
- Know your threat model
- Other goals
 - Portability
 - Performance