

More Isolation and Hardware Security

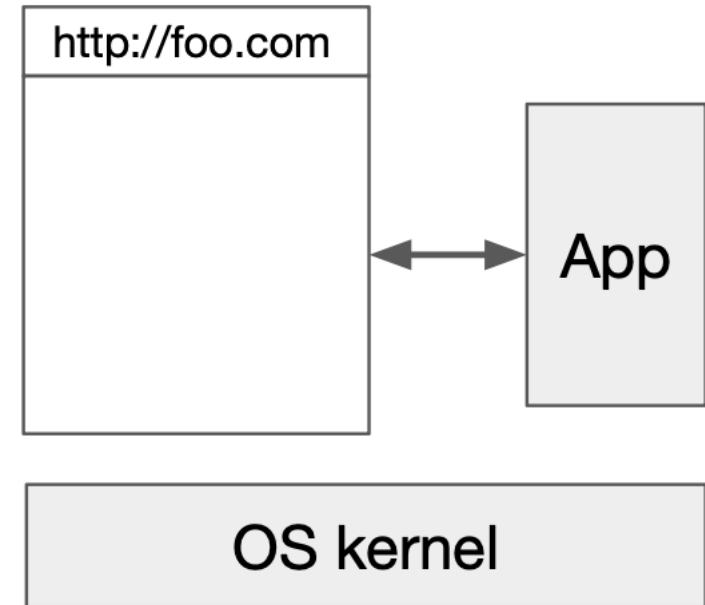
Dileepa Fernando

Overview

- Software Fault Isolation
 - Ensuring the apps are trusted
- Trusted Computing
 - Ensuring the trust in platform (OS, HW)

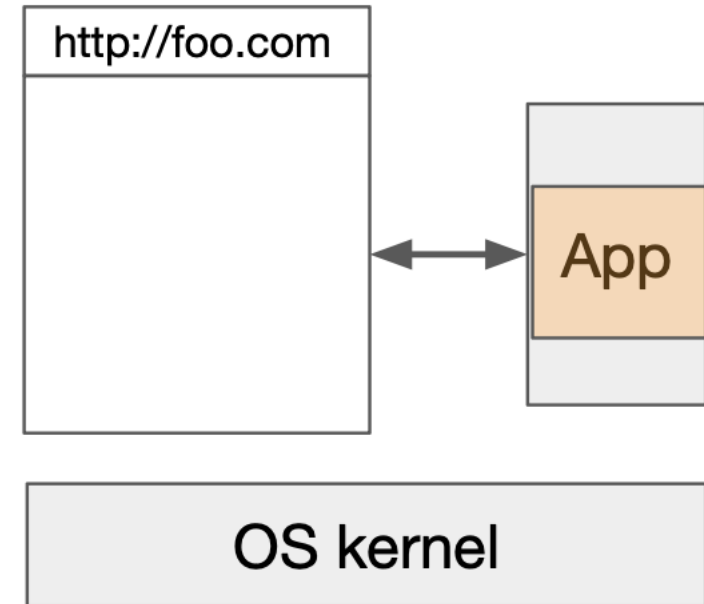
More Isolation

- Isolation So Far
 - Process
 - VM
 - Container
- Isolation enforced by manager
- Problem 1: Isolation good enough?
 - Example:
 - Native Code running on browser
 - Running as different process
 - Communicate with browser process
 - Data and Code are isolated



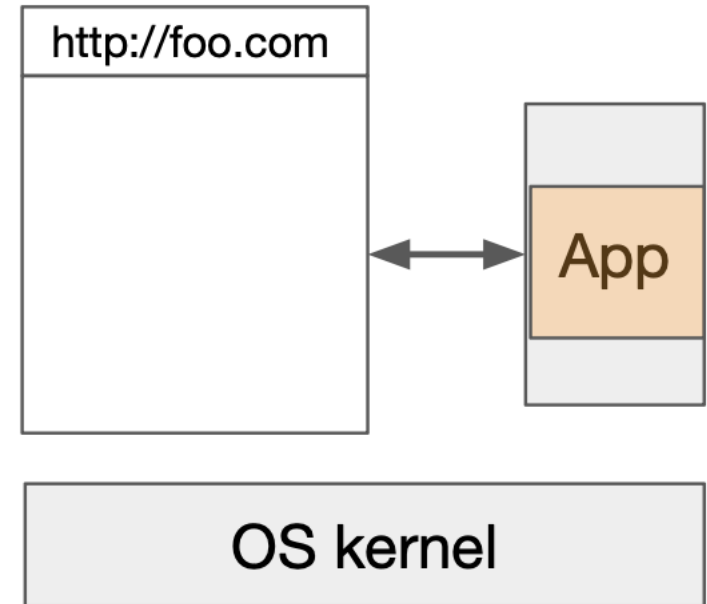
More Isolation

- Problem 1: Isolation good enough?
 - Example:
 - Native Code running on browser
 - Running as different process
 - Communicate with browser process
 - **Browser is Paranoid**
 - Does not Trust app
 - Does not Trust OS
 - Implement own isolation
 - Idea: Software Fault Isolation




More Isolation

- Problem 2: TCB size
 - Idea: Secure Hardware
 - Trusted Execution Environment (TEE)




Software Fault Isolation

- Confining apps inside sandbox
<https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/34913.pdf>
- Security Goal: App only access its own memory
 - Can communicate with other processes
 - Idea
 - Static Validation of App binary
 - Run time check of App binary
 - How sand box differs from container?
 - Security goal?
 - Threat Model?



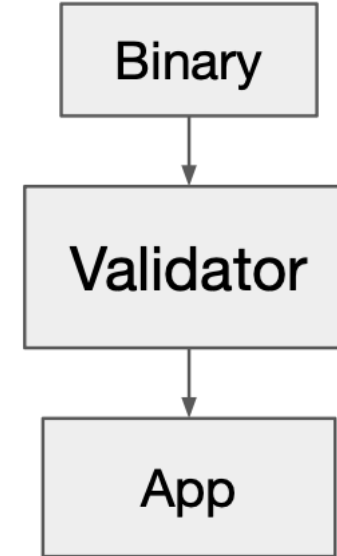
Software Fault Isolation

- Idea 1: Restricted Instructions
- Idea 2: Controlled Interaction
- Recall the goal
 - No memory access outside the app



Software Fault Isolation

- Idea 1: Restricted Instructions
 - Some instructions are safe
 - ADD, XOR
 - Allowed
 - Checked later in run time
 - Some instructions are dangerous
 - JMP
 - Insert Check
 - Some are hard to make safe
 - INT, SYSCALL
 - Disallowed



Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)

- Safe → Do nothing
- **Dangerous → Rewrite**
- Unsafe → Abort

How to rewrite?

Challenges:

- JMP addr (wrong addr)
 - JMP 0x0, JMP 0xA ok
- JMP *EAX (Not known)

f7 c7 07 00 00 00
0f 95 45 c3

test \$0x00000007, %edi
setnzb -61(%ebp)

c7 07 00 00 00 0f
95
45
c3

movl \$0x0f000000, (%edi)
xchg %ebp, %eax
inc %ebp
ret

- Any remedy?
 - Alignment (32bit)
 - No instruction consume more than 32 bit


Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)

- Safe → Do nothing
- **Dangerous → Rewrite**
- Unsafe → Abort

Challenges:

- JMP addr
 - Simple check
- JMP *EAX
- Can you jump to the middle
 - Of rewritten instruction?
 - NO
 - How RET is handled?




Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)

- Safe → Do nothing
- **Dangerous → Rewrite**
- Unsafe → Abort

More Challenges:

- Out of range access
 - JMP addr
 - JMP *EAX
- Predefine range (256MB)
- Start from 0
- Similar check before jump
 - How to?




Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)

- Safe → Do nothing
- **Dangerous → Rewrite**
- Unsafe → Abort

More Challenges:

- Out of range access
 - JMP addr
 - JMP *EAX
- Predefine range (256MB)
- Start from 0
- Similar check before jump
 - How to?



Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)
 - Safe → Do nothing
 - Dangerous → Rewrite
 - Unsafe → Abort

Any issue with the rewriting?

- Jump to forced rewrite locations
- May not be intended by a legitimate developer
- Better to just detect wrong jumps
- Use **segmentation**

Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)

- Safe → Do nothing
- **Dangerous → Rewrite**
- Unsafe → Abort

- Segment
 - Base
 - Length

- Ex: Base=0xa0000000
Length=0x1000
Addr=0xff → 0xa00000ff

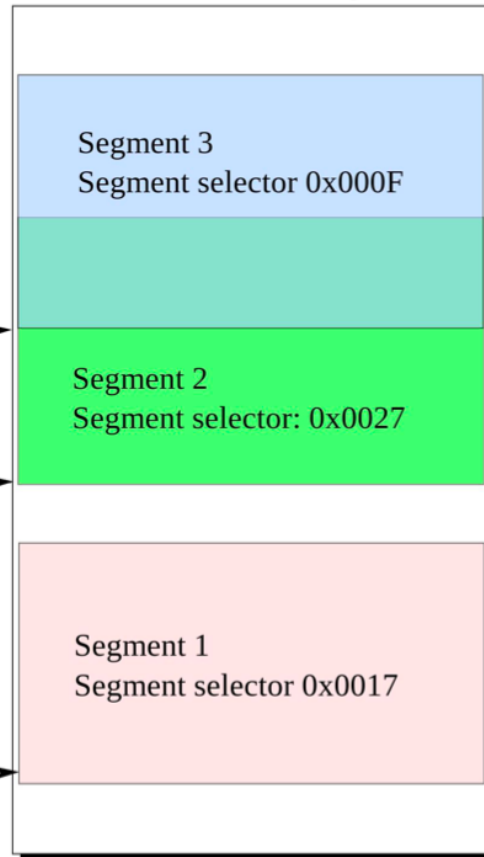
CS

GS

Local Descriptor Table (LDT)

5			
4	0x21430	0xC000	•
3			
2	0x0CEF0	0xA300	•
1	0x28C00	0xFC00	•
0			
Linear base address (BASE)		Segment size (LIMIT)	

Main memory



Software Fault Isolation

- Idea 1: Restricted Instructions (Validation)

- Safe → Do nothing
- **Dangerous → Rewrite**
- Unsafe → Abort

- Segment
 - Base
 - Length

- New app loads → Create new segment
 - Length = 256MB
 - Jump out
- Recall Segment part in IDT entry

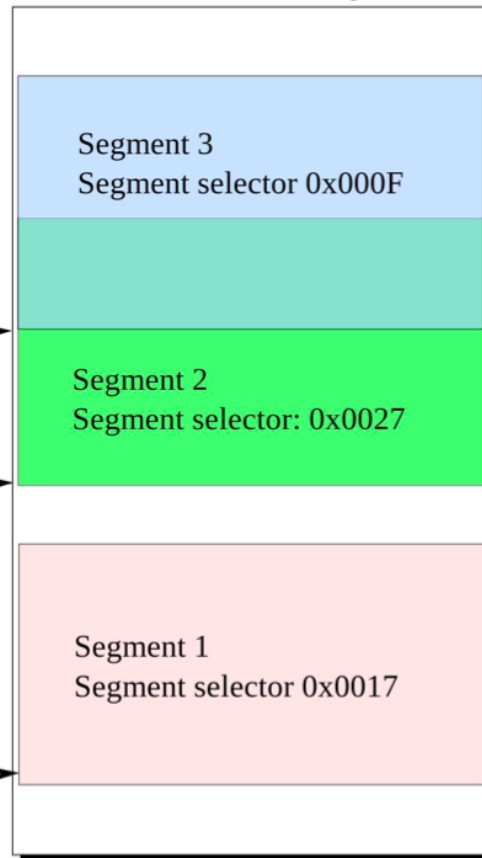
CS

GS

Local Descriptor Table (LDT)

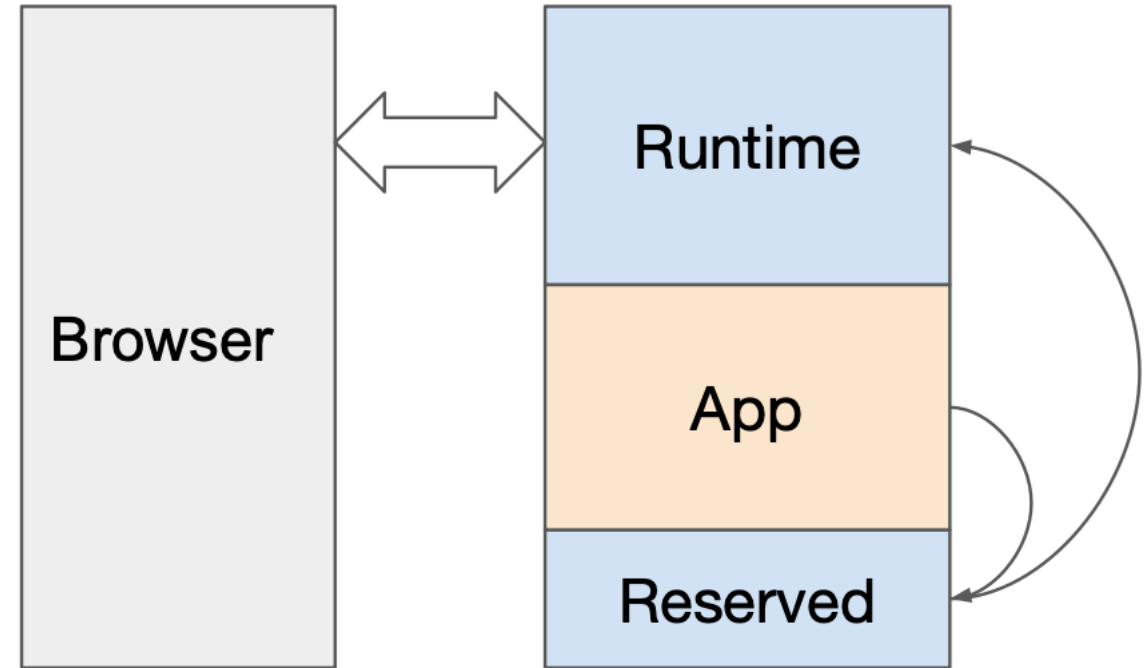
5			
4	0x21430	0xC000	•
3			
2	0x0CEF0	0xA300	•
1	0x28C00	0xFC00	•
0			
Linear base address (BASE)		Segment size (LIMIT)	

Main memory



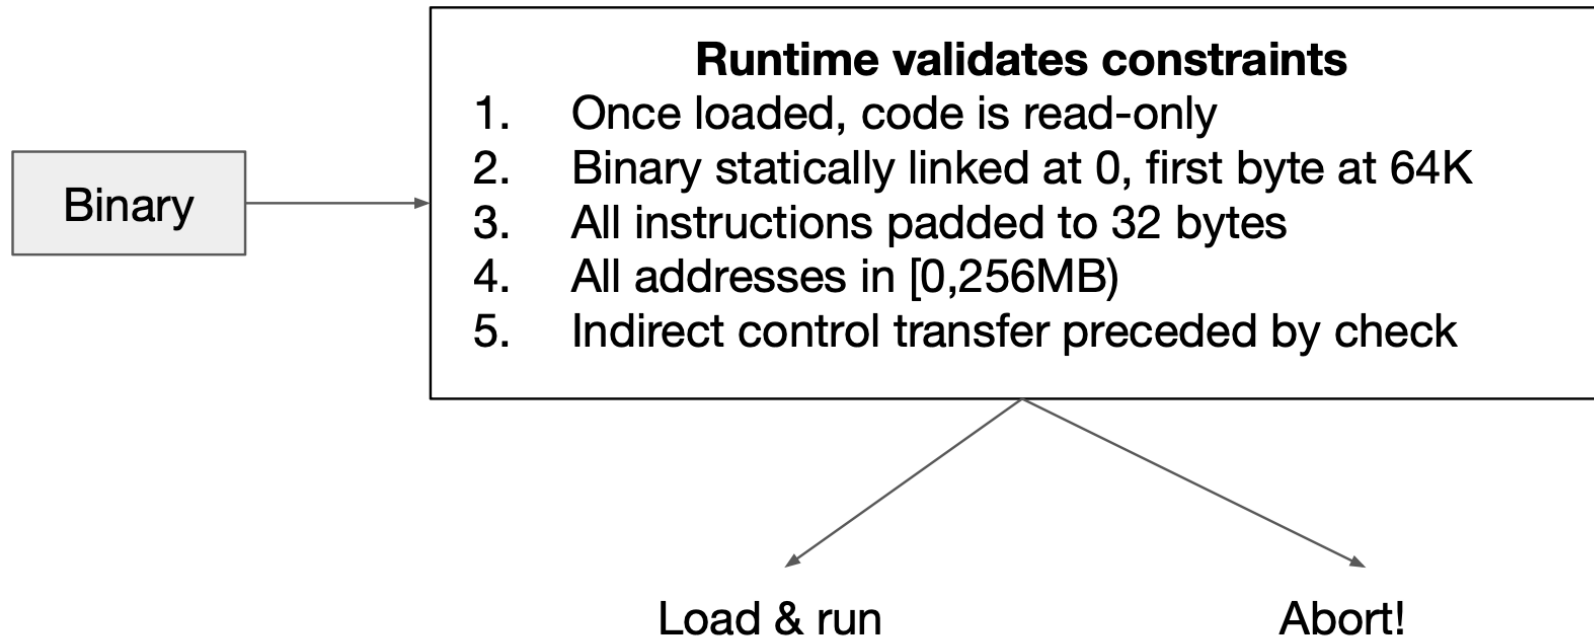
Software Fault Isolation

- Idea 2: Controlled Interaction
 - Send things to outside
 - Get things from outside
 - How?
 - Apps jump to reserved code
 - Reserved code jump to runtime (safe)
 - Runtime manages IPC with browser
 - Can the app directly jump to runtime?
 - Analogy: Process and kernel



Software Fault Isolation

- Put together



Software Fault Isolation

- Good performance
- Good Isolation


Run #	Native Client	Linux Executable
1	143.2	142.9
2	143.6	143.4
3	144.2	143.5
Average	143.7	143.3

Table 8: Quake performance comparison. Numbers are in frames per second.

- But
 - Static Code
 - Less efficient for system call heavy apps (Why?)

Story So Far

- Secure the **Applications**
- Isolation to the depth (OS, VMM, Container etc.)
- TCB = OS + Hardware (CPU rings, Interrupts, MMU)
 - But how to ensure the **trust** of OS and hardware
- What is **trust**?
 - Secure?(CIA)
 - Verified?
 - Manageable?




Hardware Security

- Security Goals: Application security
- Can we trust any OS? Can Trust a **baseline** OS
- Can we trust any Hardware? Can trust a **baseline** hardware
- Secure and Verifiable?
 - Verify the current OS/Hw with the baseline
 - But cannot always use a trusted OS baseline

Hardware Security

- Revised Security Goal: Application security Despite Malicious OS




- Detect If Cannot Prevent
 - Detect Malicious OS (TPM)
- Minimize TCB and Isolation of Secure apps
 - Trusted OS + Trusted Hw (ARM TrustZone)
 - Trusted Hw Only (Intel SGX)



Hardware Security

- **Detect malicious OS**

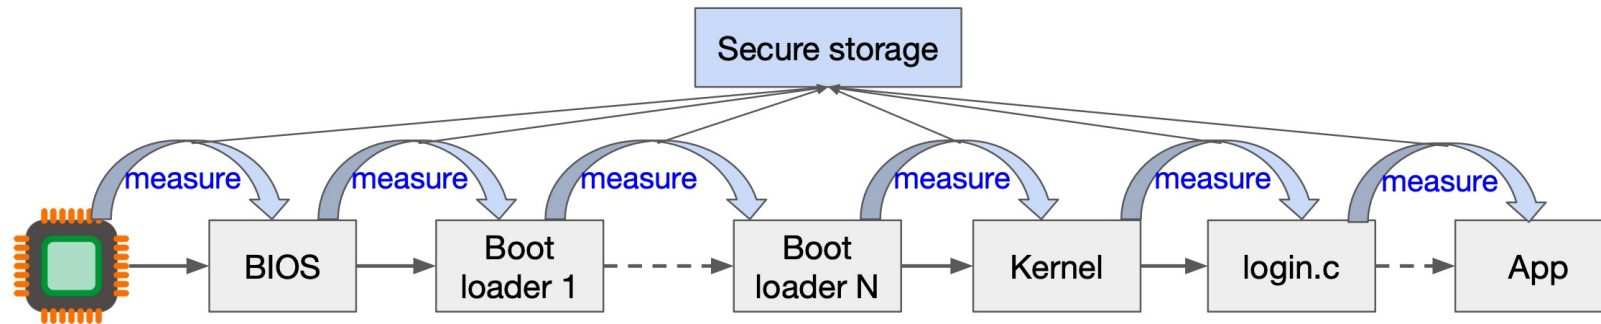
- Boot Loader N Verifies Kernel
- What if Boot Loader N is malicious?
- Verify recursively until hardware left to trust
- Remember digital certificates?



Hardware Security

- **Detect malicious OS**

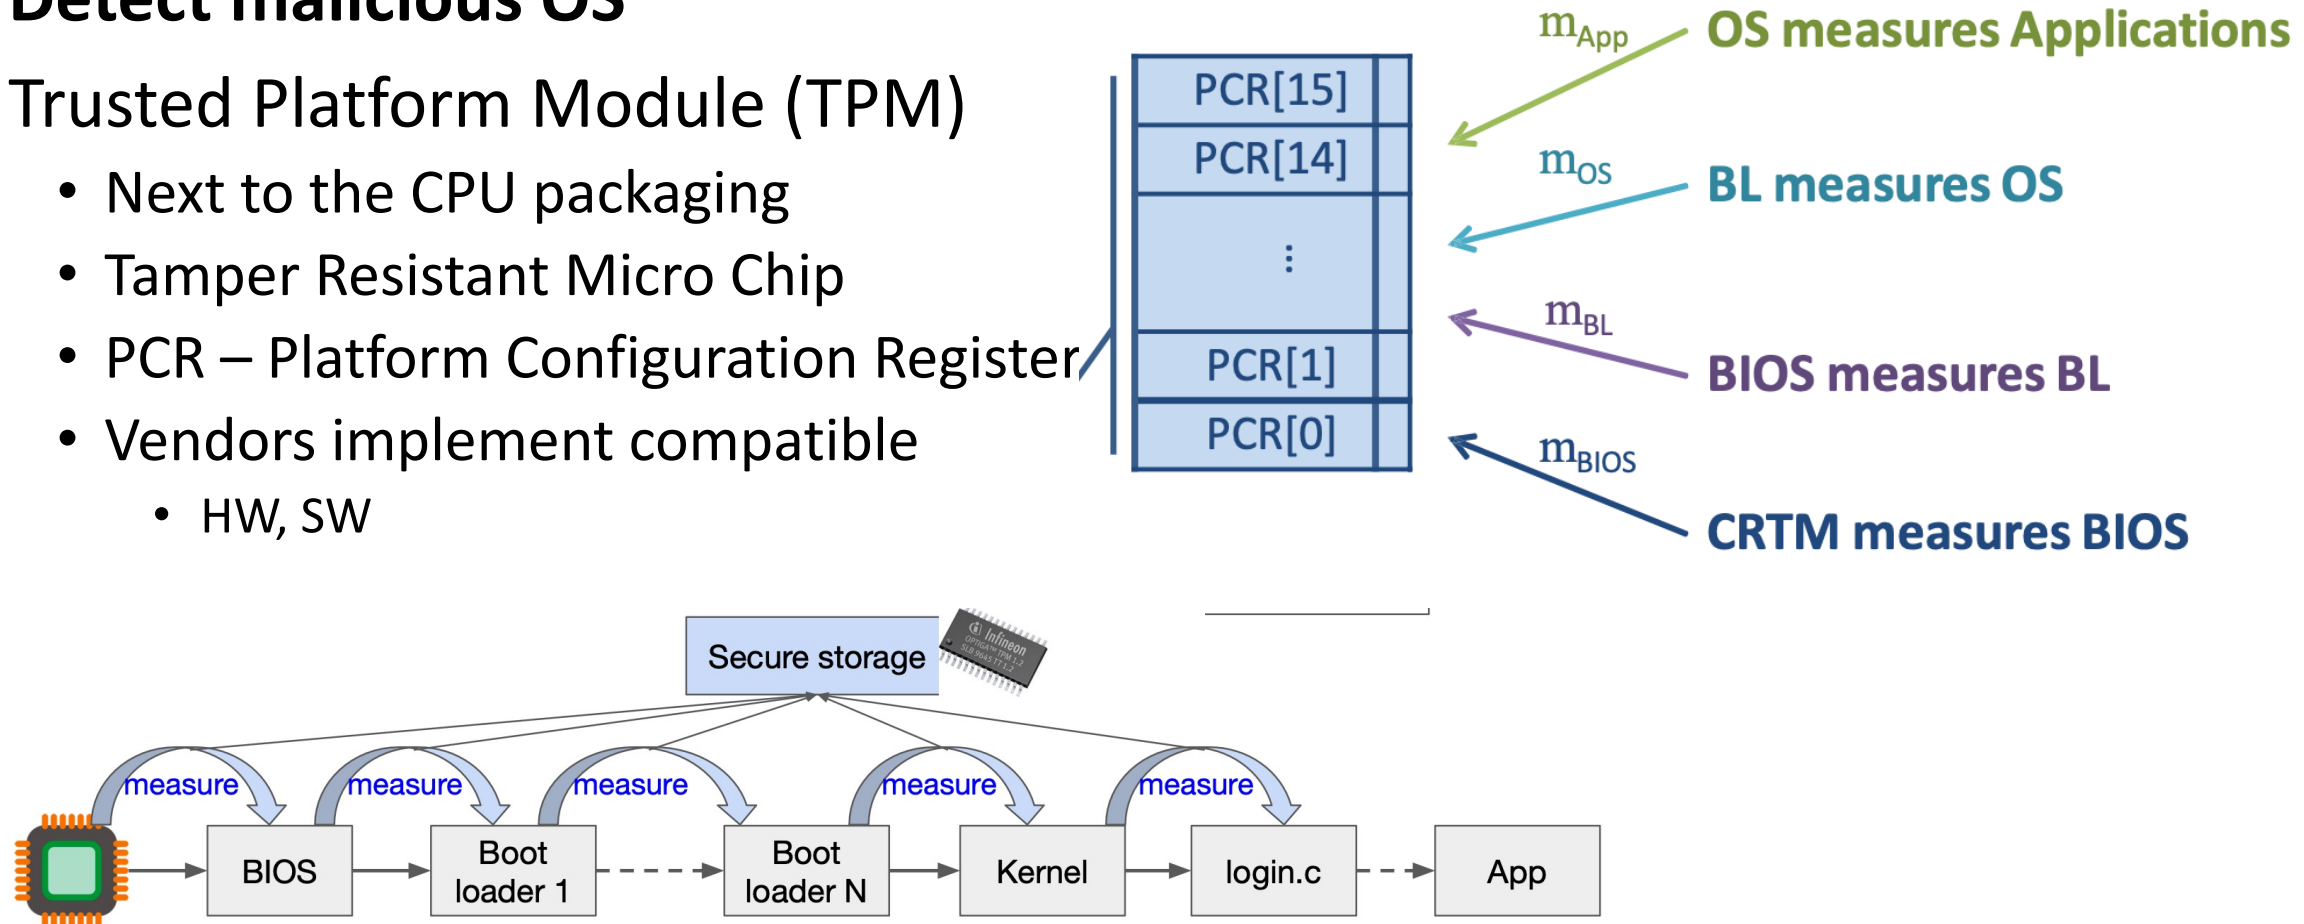
- You need to trust something
- Hardware is **root of trust**
 - Hardware measures BIOS (Core Root of Trust Management - CRTM)
 - Verified BIOS is trusted →
 - Boot Loader 1, verification is trusted →
 -Kernel,...,App verification are trusted (**chain of trust**)



Hardware Security

- **Detect malicious OS**
- **Trusted Platform Module (TPM)**
 - Next to the CPU packaging
 - Tamper Resistant Micro Chip
 - PCR – Platform Configuration Register
 - Vendors implement compatible
 - HW, SW

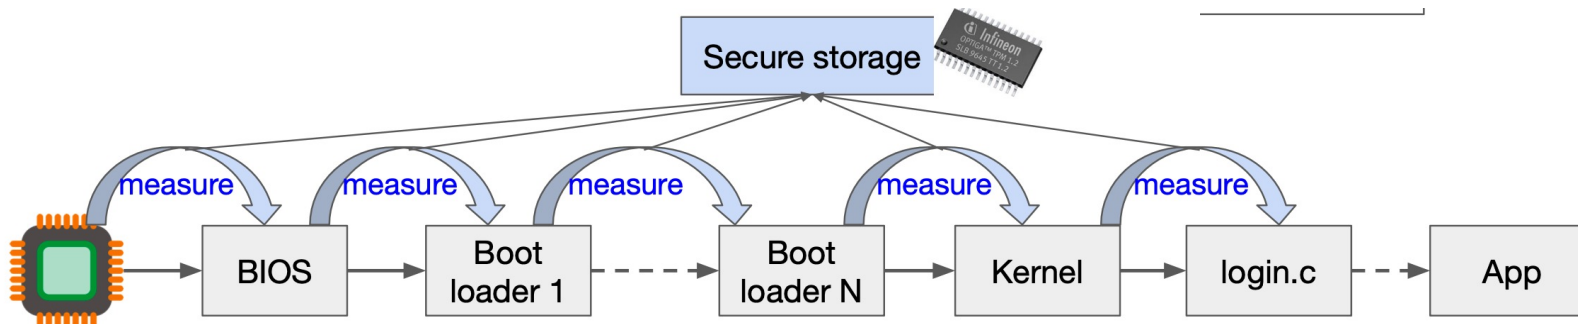
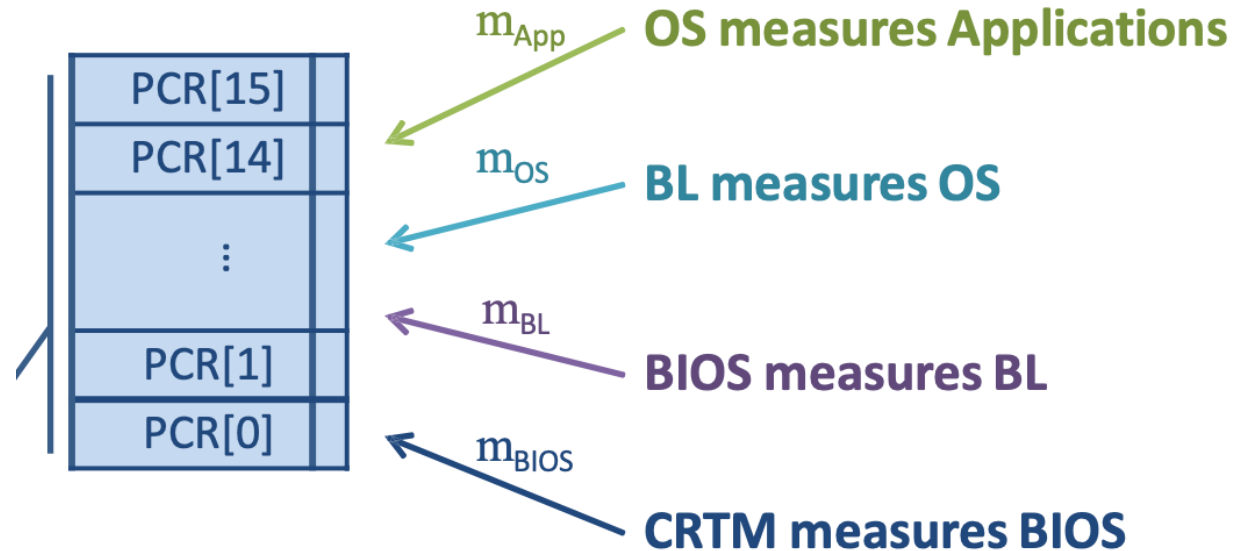
Measurement



Hardware Security

- **Detect malicious OS**
- **Trusted Platform Module (TPM)**
 - Initializes at OS installation
 - Saving Measurements in TPM
 - Verifies with every reboot
 - With TPM saved values



Measurement



Hardware Security (Initialization)

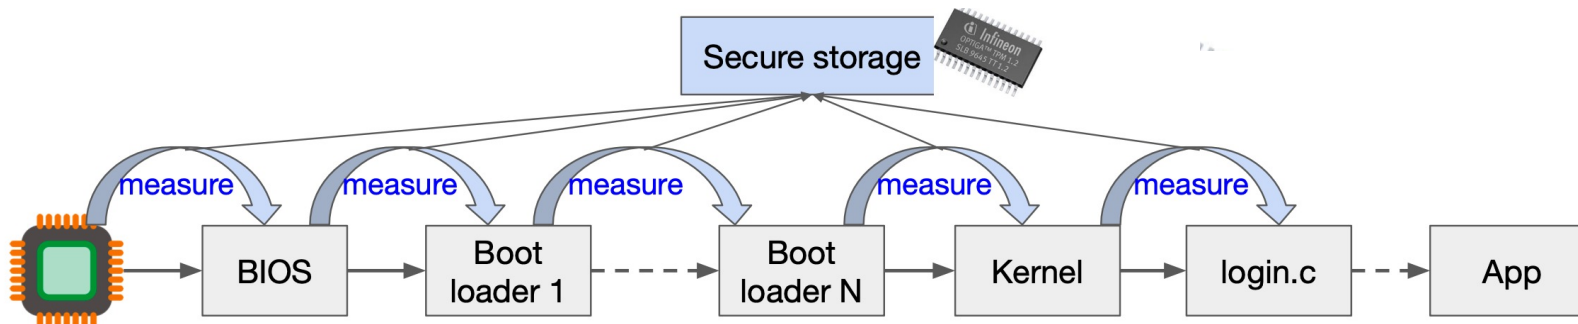
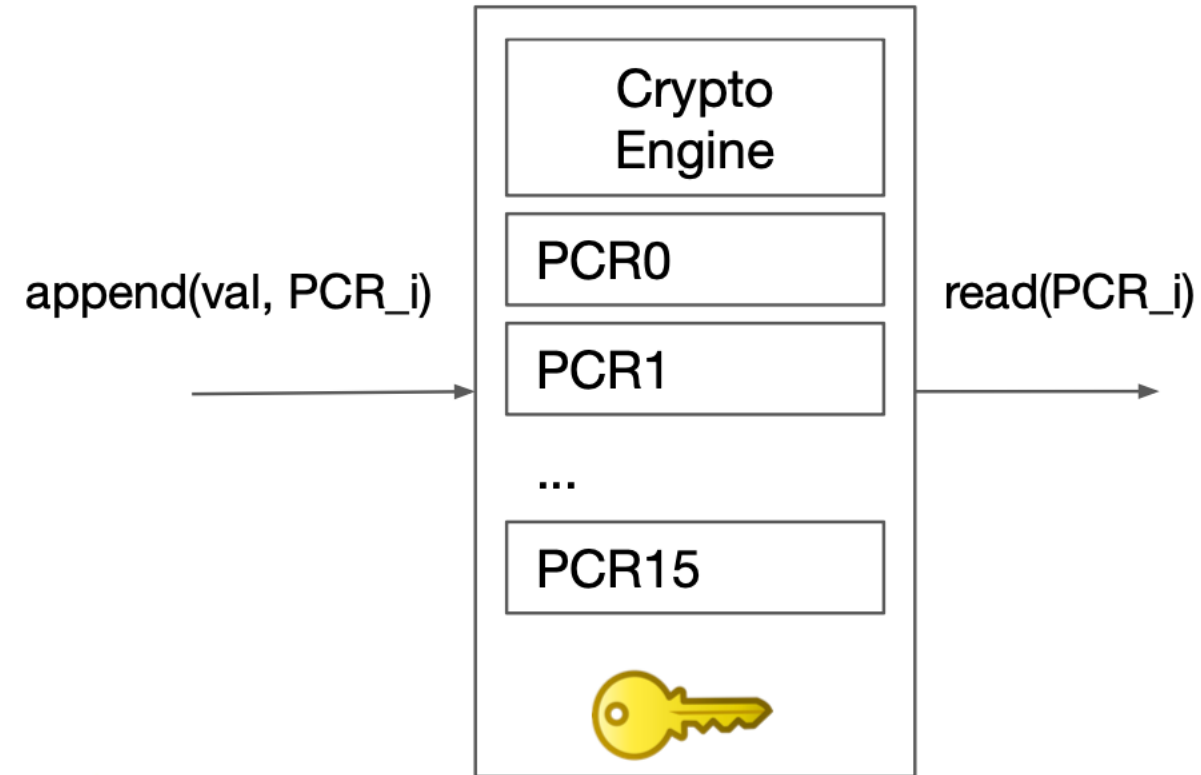
- **Detect malicious OS**
- Trusted Platform Module (TPM)
 - CRTM computes $m_{\text{BIOS}} = H(\text{BIOS})$
 - CPU call `TPM.append(m_{BIOS} , PCR_0)`
 - SET $\text{PCR}_0 = H(m_{\text{BIOS}} \parallel \text{CRTM})$

Measurement



Hardware Security (Initialization)

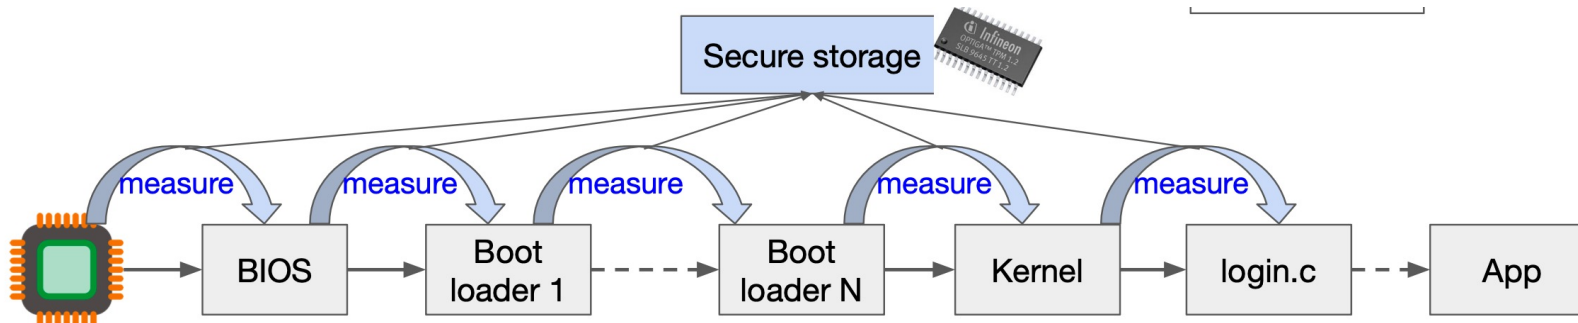
- **Detect malicious OS**
- Trusted Platform Module (TPM)
 - $m_{BL1} = H(BL_1)$, $m_{BL2} = H(BL_2), \dots$
 - $TPM.append(m_i, PCR_{i-1})$ (In order)
 - $SET PCR_1 = H(m_{BL1} || PCR_0)$
 - Follow the pattern
 - Why other PCR registers?



Hardware Security (Verification)


- **Detect malicious OS**
- Trusted Platform Module (TPM)
 - Repeated (read, append, hash) for each boot stage
 - Compare to PCR register values
 - <https://community.infineon.com/t5/Blogs/Storing-and-reporting-system-measurements-with-TPM/ba-p/443590#:~:text=PCRs%20are%20registers%20in%20TPM,storing%2020%20bytes%20of%20data>

Why does UEFI prevents dual boot?



Hardware Security

- **Detect malicious OS**
- Trusted Platform Module (TPM)
 - Remote Attestation
 - A wants to convince B
 - That it runs correct OS
- Assumptions
 - B knows correct hashes
 - B knows TPM's public keys




Hardware Security

- **Detect malicious OS**
- Trusted Platform Module (TPM)
 - Advantages
 - Very Cheap
 - Good enough for many embedded systems
 - Introduced hardware security to the main stream
 - Disadvantages
 - Slow
 - Does not guarantee OS is trusted
 - Only that it is not modified
 - Not flexible for OS updates



Hardware Security

- **Detect malicious OS**
- Trusted Platform Module (TPM)
 - Run off-chip
 - Limited memory/functions
- We want
 - More powerful/Flexible
 - Extending existing chips
 - Memory even OS cannot access
 - More privileged than OS
 - Isolated

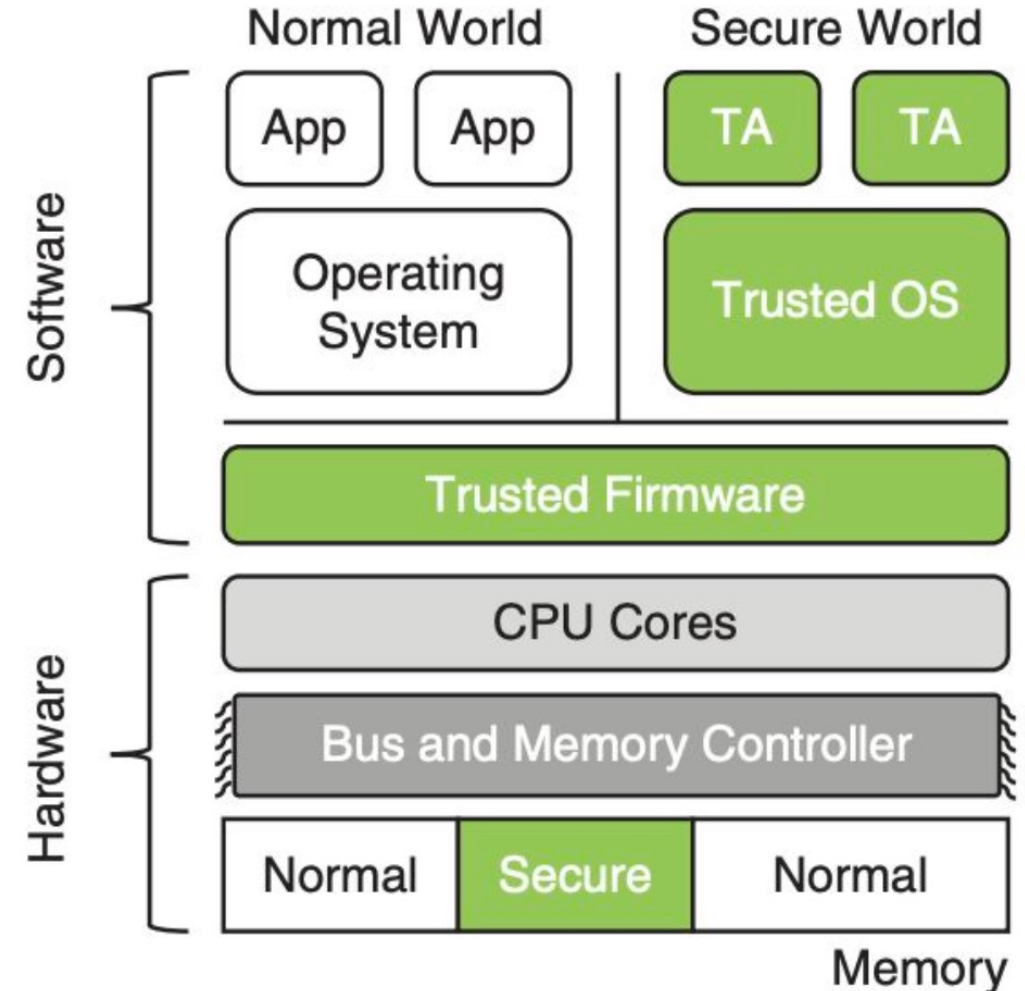


Memory (Ring 0-3)

Memory (Ring -1)

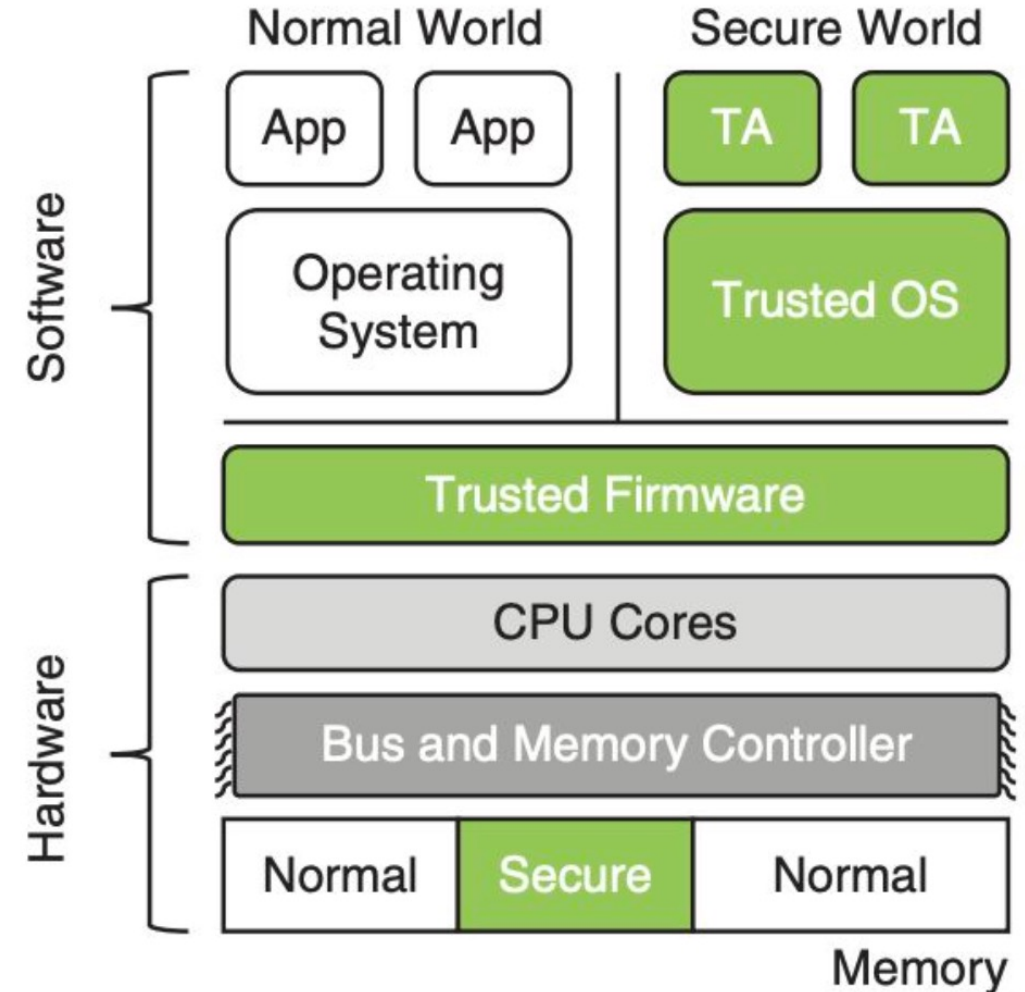
Hardware Security

- **Minimize TCB and Isolate secure apps**
- ARM's TrustZone
 - Application **Isolated** with new privilege mode
 - Non Secure Bit → 0:Secure, 1:Non-secure
 - Use Secure Memory (Physically Isolated)
 - Enforced by trusted firmware
- **Minimal TCB**
 - Trusted OS
 - Only security essential features



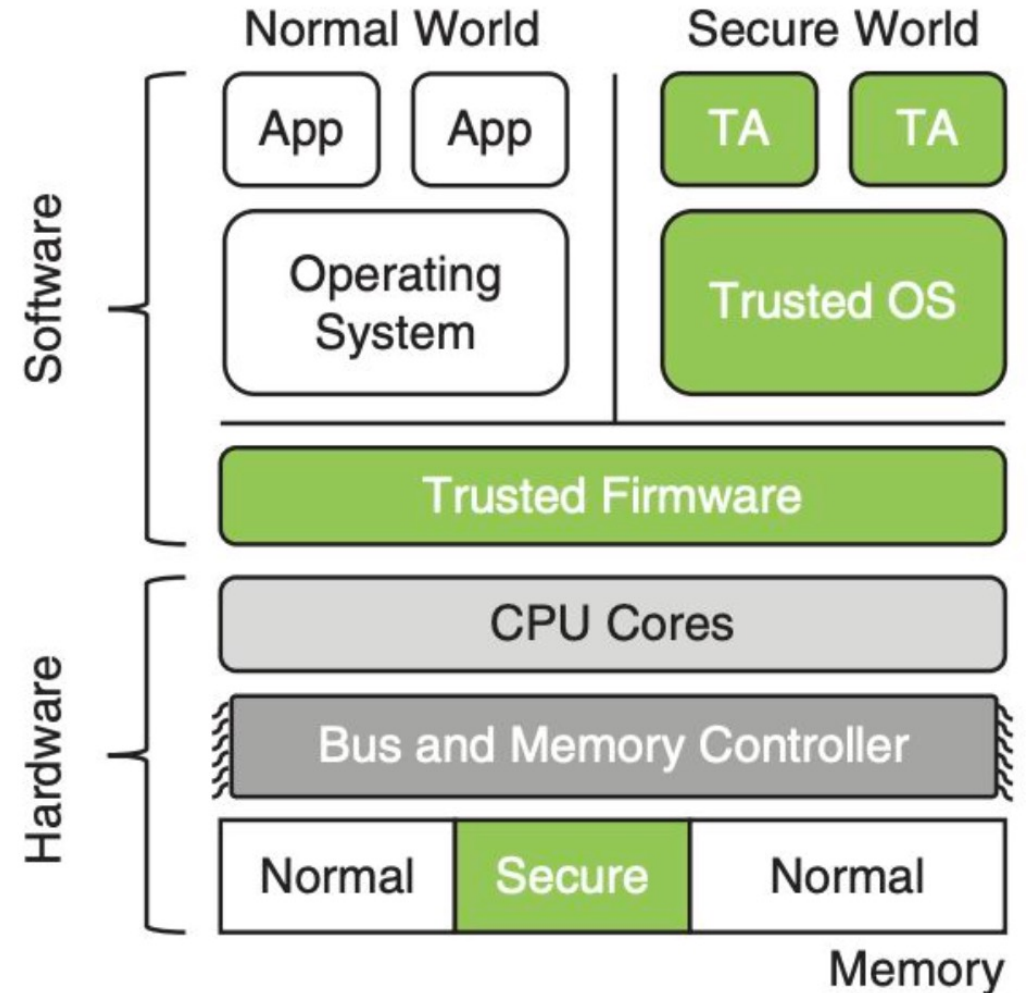
Hardware Security

- **Minimize TCB and Isolate secure apps**
- ARM's TrustZone
 - Normal world – Non sensitive app
 - Secure world – Bank app
- Trusted Firmware
 - Implements isolation (Like Kernel)
 - Set up during boot
 - Run in NS=0 mode (Monitoring Mode)
 - Secure Transition (SMC)




Hardware Security

- **Minimize TCB and Isolate secure apps**
- Application Security
 - Trust verification
 - Authenticated Boot
 - Remote Attestation
 - TCB reduced
 - Trusted OS/Firmware
 - Still big
 - Physical memory attacks?
 - Secure memory is not encrypted




Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Application **Isolated**
 - Enclave mode vs Non-enclave mode
 - Enclave mode is highest privilege
 - Processor Reserved Memory (PRM)
 - Encrypted
 - Enforced by SGX Micro Code
 - **Minimal TCB**
 - SGX Micro Code
 - Memory Encryption Engine
 - CPU does secure memory management



Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Memory Isolation
 - PRM dedicated for enclaves
 - Enclave Page Cache (EPC)
 - Allocatable pages
 - Allocated pages **encrypted**
 - Enclave Page Cache Map (Page Table)
 - One Entry for each page
 - **CPU manages** (not OS)
 - Only CPU can see the mem. layout
 - Only CPU can edit page table



Hardware Security

- **Minimize TCB and Isolate secure apps**

- Intel's Secure Guard Extension (SGX)



- Memory Isolation

- Role of OS

- Keep track of EPC/normal spaces
 - Normal process → Normal page
 - Enclave process → EPC
 - EADD instruction
 - CPU update the EPCM

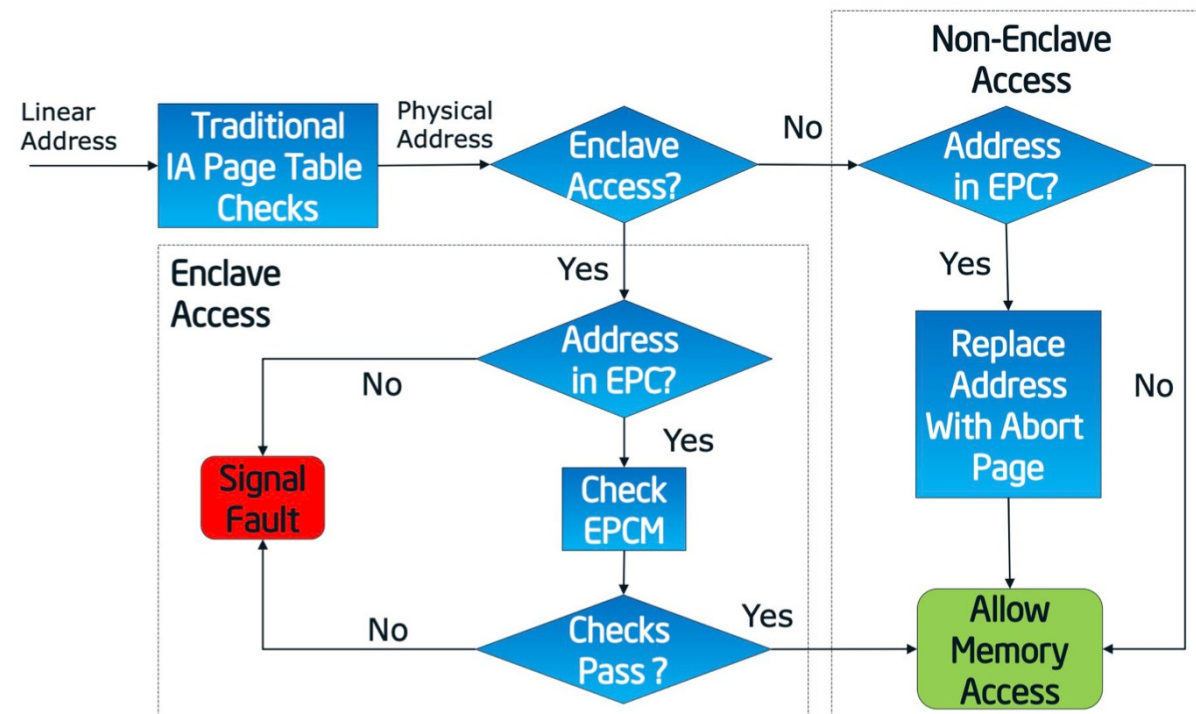
- Transition between modes

- EENTER, ERESUME
 - EXIT, AEX




Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Memory Translation
 - MMU translation
 - Based on page tables
 - Non enclave page table – OS managed
 - Enclave page table – CPU managed
 - Faults
 - Signal fault – CPU error
 - Abort page – OS error
 - OS never sees EPC layout



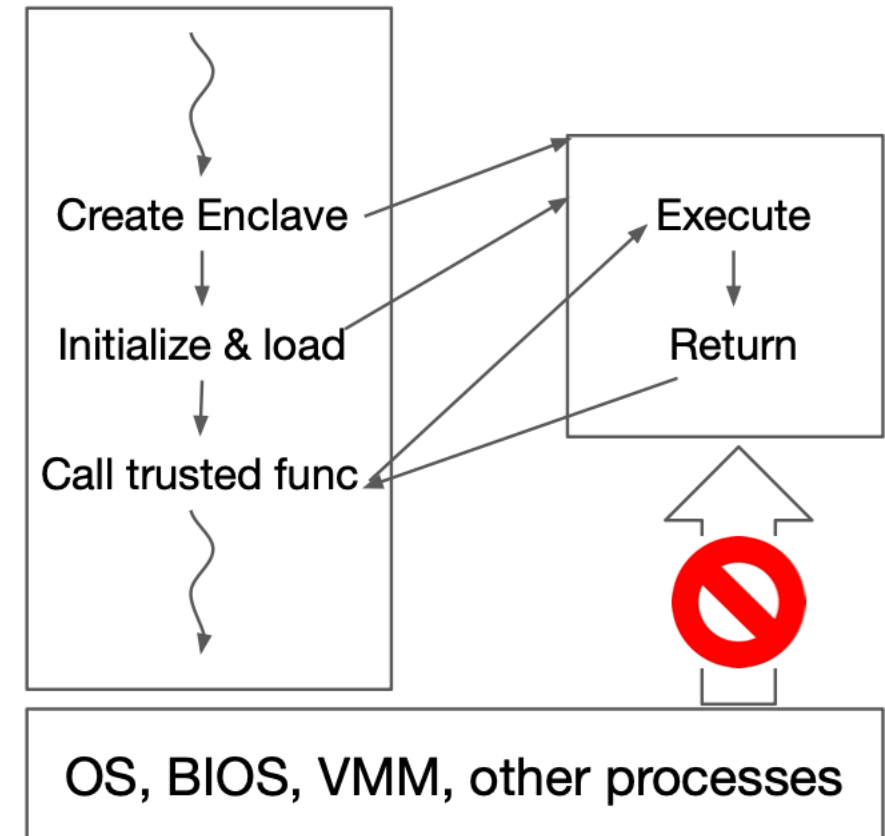
Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Programming Model
 - Application has two parts
 - Trusted (Enclave)
 - Non-trusted
 - Create, Init, Run, Exit
 - Instructions
 - Can OS read from enclave?
 - MOV <addr>, EAX
 - No.
 - Can Enclave access untrusted memory?
 - Not directly



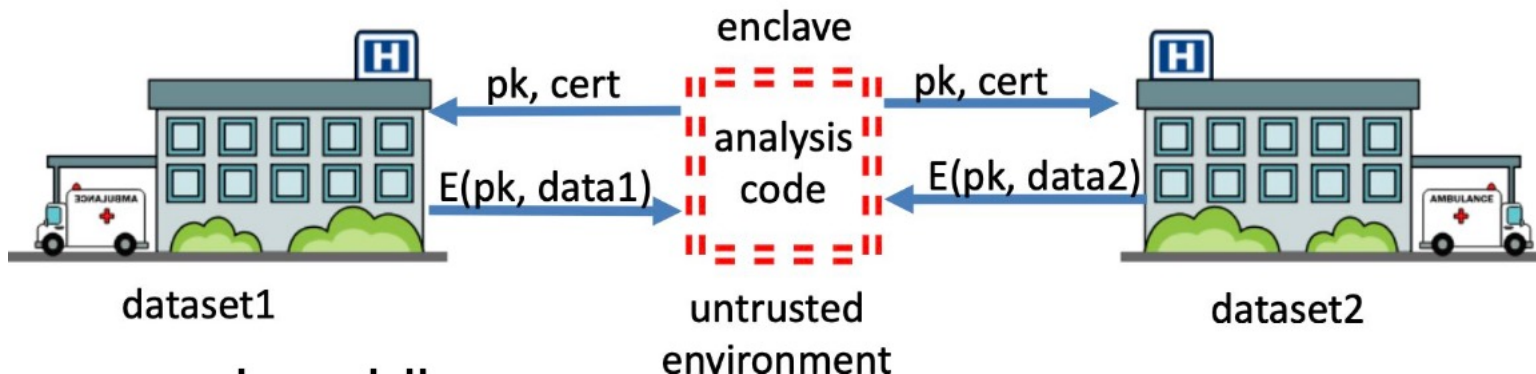
Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Programming Model
 - Multi Threading
 - Context Switch from Enclave to OS
 - CPU cleans up
 - Attestation
 - CPU measures Enclave (Code and data)
 - Signs
 - Remote party verifies



Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Example
 - Hospital encrypts a dataset (Source Code Public)
 - Runs analysis on a server (SGX supported)
 - Communicate using public key
 - Private key stored in Memory encryption engine



Hardware Security

- **Minimize TCB and Isolate secure apps**
- Intel's Secure Guard Extension (SGX)
 - Advantages
 - Small TCB
 - Commercially Available
 - Disadvantages
 - Side channel attacks
 - Physical attacks?
 - Encrypted
 - Can someone physically access memory encryption engine?

Hardware Security

Summary:

- OS and HW, TCB
- How to Trust OS and HW
- Verify Known Firmware and OS (TPM by Trusted Computing Group)
- Extended with hardware enforced isolation (ARM Trustzone)
 - Trusted OS/Firmware
- Extended with memory encryption (Intel SGX)
 - Trusted CPU/Firmware